



Chapter 4

Threads & Concurrency

Da-Wei Chang

CSIE.NCKU

Source: Abraham Silberschatz, Peter B. Galvin, and Greg Gagne, "Operating System Concepts", 10th Edition, Wiley.

Outline

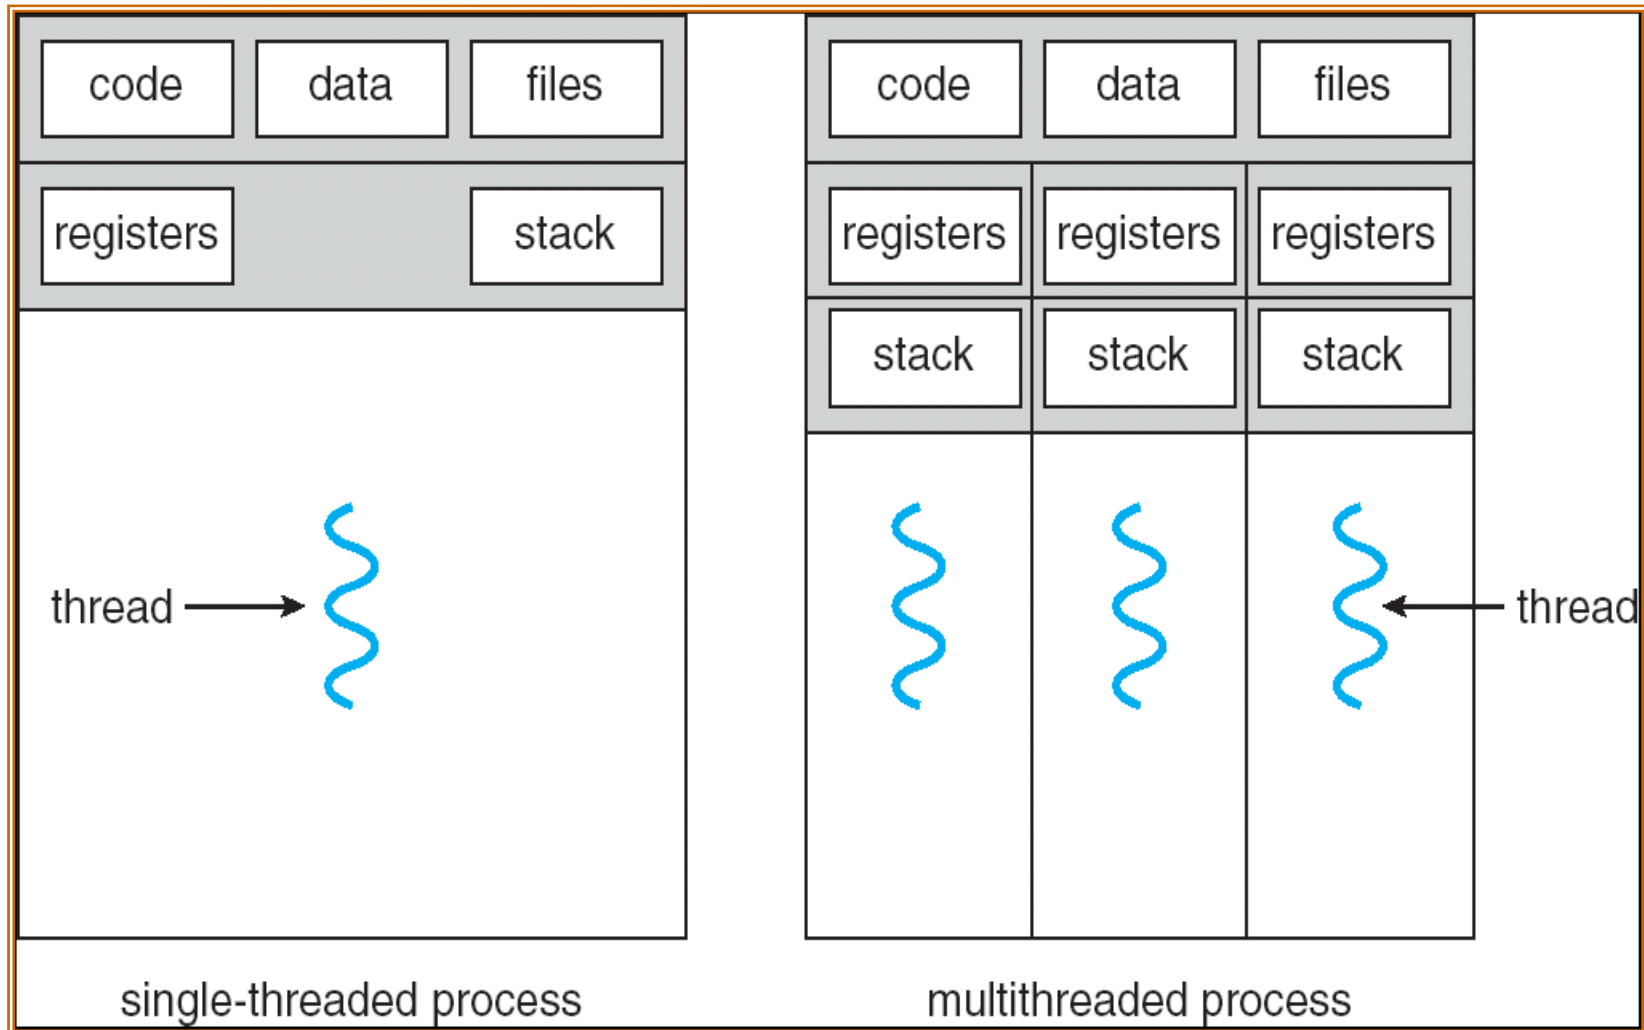


- Overview
- Multithreading Models
- Thread Libraries
- Threading Issues
- Operating-System Examples

Overview

- Most operating systems support multi-threaded processes
- A thread
 - A basic unit of CPU utilization
 - Comprises
 - Thread id
 - Program counter
 - Register set
 - Stack
- A multi-threaded process can perform more than one task at a time

Single and Multithreaded Processes



Motivation

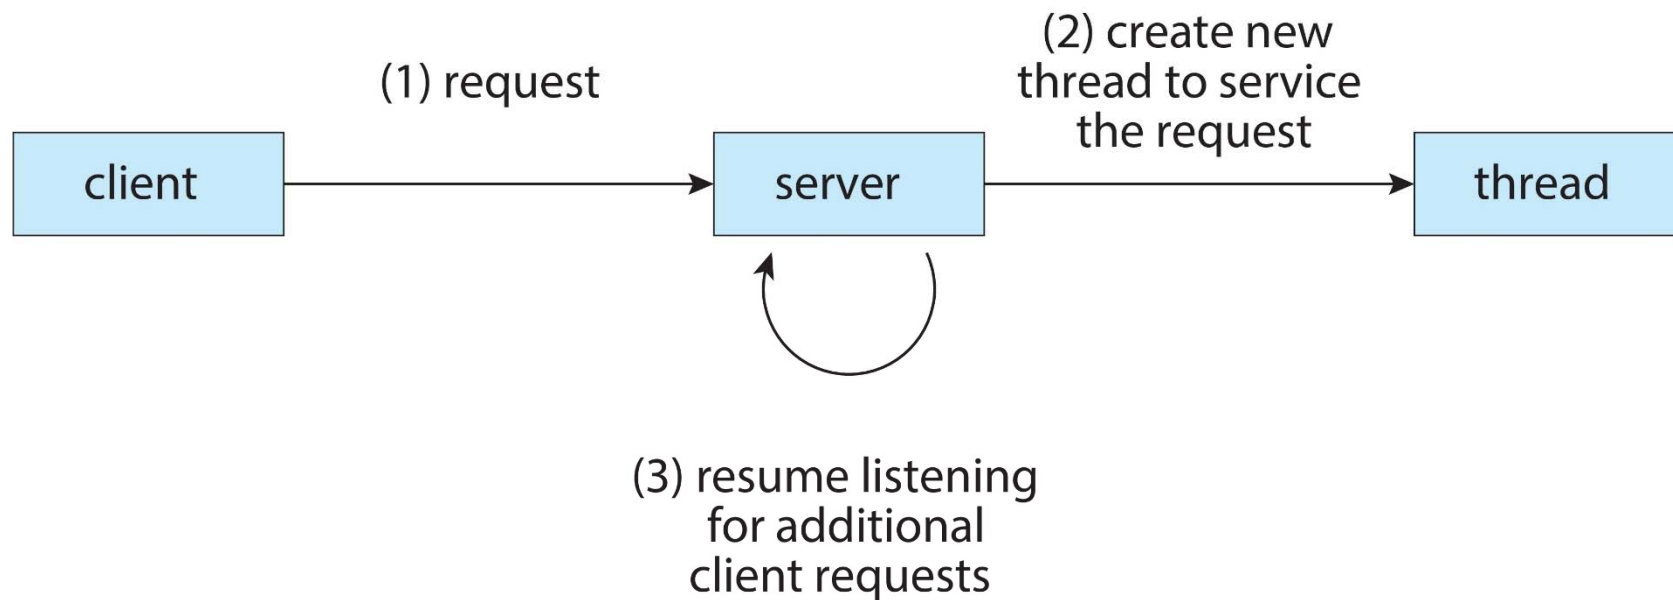
- A lot of software packages are multi-threaded
 - Web browser
 - One thread displays images/text
 - Another retrieves data from the network
 - Word processor
 - Display graphics
 - Respond to keystrokes
 - Perform spelling and grammar checking in the background
 - Web server
 - May have one thread for each request

Motivation

- If a web server is single-threaded
 - Serve only one client at a time
 - If it processes the requests one-by-one → a long waiting time
- Multi-process solution
 - Common before threads become popular
 - Process creation is time consuming and resource intensive
- Multi-threaded solution
 - More efficient
 - Threads are more lightweight than processes
- A multi-threaded web server may
 - Use a thread for listening for client requests
 - Create a thread for handling each request

** see the next slide*

Multithreaded Server Architecture



Motivation

- Threads are also important in RPC systems
 - RPC servers are multi-threaded
 - Serve each message via a separated thread
 - Java's RMI systems are also multi-threaded
- Many operating systems are multi-threaded
 - Each thread performs a specific task
 - E.g., Solaris has a set of threads for interrupt handling
 - E.g., Linux uses kernel thread(s) for writing back memory data to disk

Benefits


- Responsiveness
 - Allows a program to continue running even if part of it is blocked
 - E.g., a web browser allows user interaction while loading the text/image
- Utilization of MP Architectures
 - Threads can run in parallel on different processors
 - Allows a multithreaded application to run on top of multiple processors
- Resource Sharing
 - Threads share memory & resources
 - Lightweight communication through memory sharing
- Economy
 - More economic to create/switch threads
 - In Solaris, process creation is 32x slower, process switching is 5x slower
- The former two also apply to multi-process architectures, while the latter two are more specific to multi-threading.

Multiprocessor Programming


- **Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
- ***Parallelism*** implies a system can perform more than one task simultaneously
- ***Concurrency*** supports more than one task making progress
 - Single processor/core, scheduler providing concurrency

Concurrency vs. Parallelism

- **Concurrent execution on single-core system:**





- **Parallelism on a multi-core system:**



Multiprocessor Programming


- Types of parallelism
 - **Data parallelism** – distributes subsets of the same data across multiple cores, **same operation** on each
 - **Task parallelism** – distributing threads across cores, each thread performing **unique operation**



Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
 - S is serial portion
 - N processing cores
- $$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$
- An example
 - if application is 75% parallel and 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
 - As N approaches infinity, speedup approaches $1 / S$

Amdahl's Law



Thread Types

- User threads
- Kernel threads

User Threads


- User-level threads
- Thread management done by **thread library**
- Common thread libraries (*not always user-level implementations, described later*):
 - POSIX Pthreads
 - POSIX: Portable Operating System Interface
 - Windows threads
 - Java threads

Kernel Threads



- Supported by the Kernel
- Examples
 - Windows
 - Solaris
 - Linux
 - Tru64 UNIX (formerly Digital UNIX)
 - Mac OS X

User and Kernel Threads



Multithreading Models




- Many-to-One
- One-to-One
- Many-to-Many

Many-to-One

- Many user-level threads mapped to single kernel thread
- Thread management is done by the **thread library** in the **user space**
 - Efficient on thread management
 - If one thread makes a blocking system call, the other threads mapping to the same kernel thread will block
 - Multiple threads are **unable** to run in parallel on multiprocessors
 - since the kernel only sees a single thread (for these user threads)
- Few systems currently use this model
- Examples
 - Solaris Green Threads
 - GNU Portable Threads


Many-to-One Model



One-to-One

- Each user-level thread maps to a kernel thread
- Allows another thread to run when a thread makes a blocking system call
- Multiple threads can run in parallel on multiprocessors
- The overhead of the kernel threads may burden the performance of the applications
 - Therefore, sometimes you can not create as many threads as necessary
- Examples
 - Windows
 - Linux
 - Solaris 9 and later


One-to-one Model



Many-to-Many Model

- Allows M user level threads to be mapped to N kernel threads, where $N \leq M$
- Allows the operating system to create a sufficient number of kernel threads
- Allows applications to create as many user threads as necessary
- The kernel threads can run in parallel on multiprocessors
- A blocking system call does not block the entire process
- Examples
 - Solaris prior to version 9
 - Windows NT/2000 with the *ThreadFiber* package

Many-to-Many Model




Two-level Model



- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier

Two-level Model



Thread Libraries

- Provides an API for creating and managing threads
- Two primary ways of implementation
 - User-level library
 - Code and data structures reside entirely in the user space
 - A library call does not result in a system call
 - Kernel-level library
 - Code and data structures reside in the kernel space
 - A library call typically results in a system call

Thread Libraries



- Common Thread Libraries
 - Pthreads
 - Provided as either a user- or kernel-level library
 - Windows threads
 - Kernel-level library
 - Java threads
 - Create and manage java threads
 - Use native threading support, typically
 - Use Windows threads in Windows
 - Use Pthreads in UNIX or Linux

Pthreads

- POSIX threads
 - A specification for thread behavior (IEEE 1003.1c)
 - Not an implementation!!!
 - OS designers can implement the specification
 - Numbers of implementations in
 - Solaris, Linux, Mac OS X, Tru64 UNIX
 - Shareware implementations in Windows
- A Pthreads example (see next two slides)
 - A thread begins with `main()`
 - `main()` creates a second thread by `pthread_create()`
 - Both threads share the global variable *sum*
 - Wait for a thread to terminate by `pthread_join()`

A Pthreads Example

```
int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* the thread */

main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */
    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        exit();
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be <= 0\n", atoi(argv[1]));
        exit();
    }
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* now wait for the thread to exit */
    pthread_join(tid, NULL);
    printf("sum = %d\n", sum);
}
```

A Pthreads Example (cont.)

```
/* The thread will begin control in this function */  
void *runner(void *param)  
{  
    int upper = atoi(param);  
    int i;  
    sum = 0;  
    if (upper > 0) {  
        for (i = 1; i <= upper; i++)  
            sum += i;  
    }  
    pthread_exit(0);  
}
```


Windows Threads



- Threads are created by `CreateThread()`
- `WaitforSingleObject()`
 - Wait for the specified thread to terminate

Windows Threads

```
int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    Param = atoi(argv[1]);
    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
```

Java Threads

- Two techniques to create Java threads
 - Create a new class deriving from the `Thread` class and override its `run()` method
 - Define a class that implements the `Runnable` interface

```
Public interface Runnable
{
    Public abstract void run();
}
```
- Threads are actually created when the `start()` method of the thread object is invoked
 - The `start()` method
 - Allocates memory and init a new thread in the JVM
 - Calls the `run()` method
- Invoke `join()` method of the thread object to wait for the thread to exit

Implicit Threading



- Growing in popularity
 - as numbers of threads increase, program correctness more difficult with explicit threads
- Creation and management of threads done by compilers and run-time libraries rather than programmers
- Examples
 - OpenMP
 - Intel Threading Building Blocks

OpenMP

- Set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that can run in parallel

#pragma omp parallel

- create as many threads as there are cores

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```

OpenMP

- Run the for loop in parallel

```
#pragma omp parallel for  
for (i = 0; i < N; i++) {  
    c[i] = a[i] + b[i];  
}
```

Threading Issues



- Semantics of **fork()** and **exec()** system calls
- Thread cancellation
- Signal handling
- Thread pools
- Thread specific data

Semantics of `fork()` and `exec()`

- Does **`fork()`** duplicate only the calling thread or all threads?
- Some UNIX systems support two versions of `fork()`
 - Duplicate all threads
 - Duplicate the caller thread only
- How about `exec()`?
 - A thread invokes `exec()` → **entire** process will be replaced, including all of the threads
- Which version of `fork()` should be used?
 - Depends on whether the `exec()` will be called
 - If it will → just duplicate the caller
 - Otherwise, duplicate all the threads

Thread Cancellation

- Thread Cancellation
 - Terminating a thread **before** it has finished
- Examples
 - Searching a database
 - If one thread finds the result, cancel the other threads
 - Web browser
 - If user press the cancel button, cancel the downloading thread
- Two general approaches for thread cancellation
 - **Asynchronous cancellation** terminates the target thread **immediately**
 - May have problems when
 - Resources have been allocated to the target thread
 - The target thread is updating a shared data
 - **Deferred cancellation** allows the target thread to **check later** if it should be cancelled
 - Allows a thread to be cancelled in a safe point (**cancellation point**)

Thread Cancellation (Cont.)

- Cancellation mode depends on *state* and *type*

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If a thread has cancellation state disabled, cancellation remains pending until the thread enables it
- Cancellation points for deferred cancellation
 - certain functions must, and certain other functions may, be cancellation points
 - For example, `open()/close()/read()/write()` are cancellation points
 - See <http://man7.org/linux/man-pages/man7/pthreads.7.html>
 - Insert a cancellation point: **`pthread_testcancel()`**

Signal Handling

- Signals are used in UNIX systems to **notify a process** that a particular **event** has occurred
- A signal may be received either synchronously or asynchronously
 - **Synchronous** signals
 - E.g., illegal memory access, divided by 0
 - Generated by the process itself
 - **Asynchronous** signals
 - E.g., terminate a process with Ctrl-C, timer expires
 - Generated by an external event

Signal Handling

- A **signal handler** is used to process signals
 - Every signal has a default handler that kernel runs when handling signal
 - User-defined signal handler can override default
 - Signal generation & handling
 1. Signal is generated by particular event
 2. Signal is delivered to a process (signal destination)
 3. Signal is handled
 - Signal destination for a multi-threaded process
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process
- *A signal can be handled only once!**

Signal Handling

- For syn signals, it should be delivered to the thread that causes the signal
- For asyn signals, it depends...
 - E.g., Ctrl-C → the signal should be sent to all the threads
- Many OSes allows a **thread** to specify which signals to accept and which signals to block
 - A signal can be delivered to the **first found** thread that accepts it
 - A signal can be handled only once!

Signal Handling

- Sending a signal to a process in UNIX
 - `kill (pid, signal)`
- Sending a signal to a thread in Pthreads
 - `pthread_kill (tid, signal)`
- Asynchronous Procedure Call (APC)
 - in Windows systems
 - Similar to asyn signals
 - Allows a thread to specify a function to be called when an event happens
 - delivered to a thread, not a process

Thread Pools

- Consider a web server that serves each request with a separated thread
 - Thread creation and termination are not free
 - No upper limit on the number of threads
- Solution: **Thread Pool**
 - Create a number of threads in a pool where they await work
- Advantages
 - Usually slightly faster to serve a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool

Thread Pools

- The number of threads in the pool can be based on
 - Number of CPUs
 - Amount of physical memory
 - Expected number of requests
- Some thread pool architectures can adjust the number of threads in the pool according to usage patterns or load

Thread-Specific Data

- Also called Thread-Local Storage (TLS)
- Allows each thread to have its own copy of data
- Thread specific data is **private** to a thread, but **shared** among the functions invoked by the thread
- Most thread libraries support this feature
- Two pthreads functions for thread specific data
 - pthread_setspecific()
 - pthread_getspecific()

Thread Examples



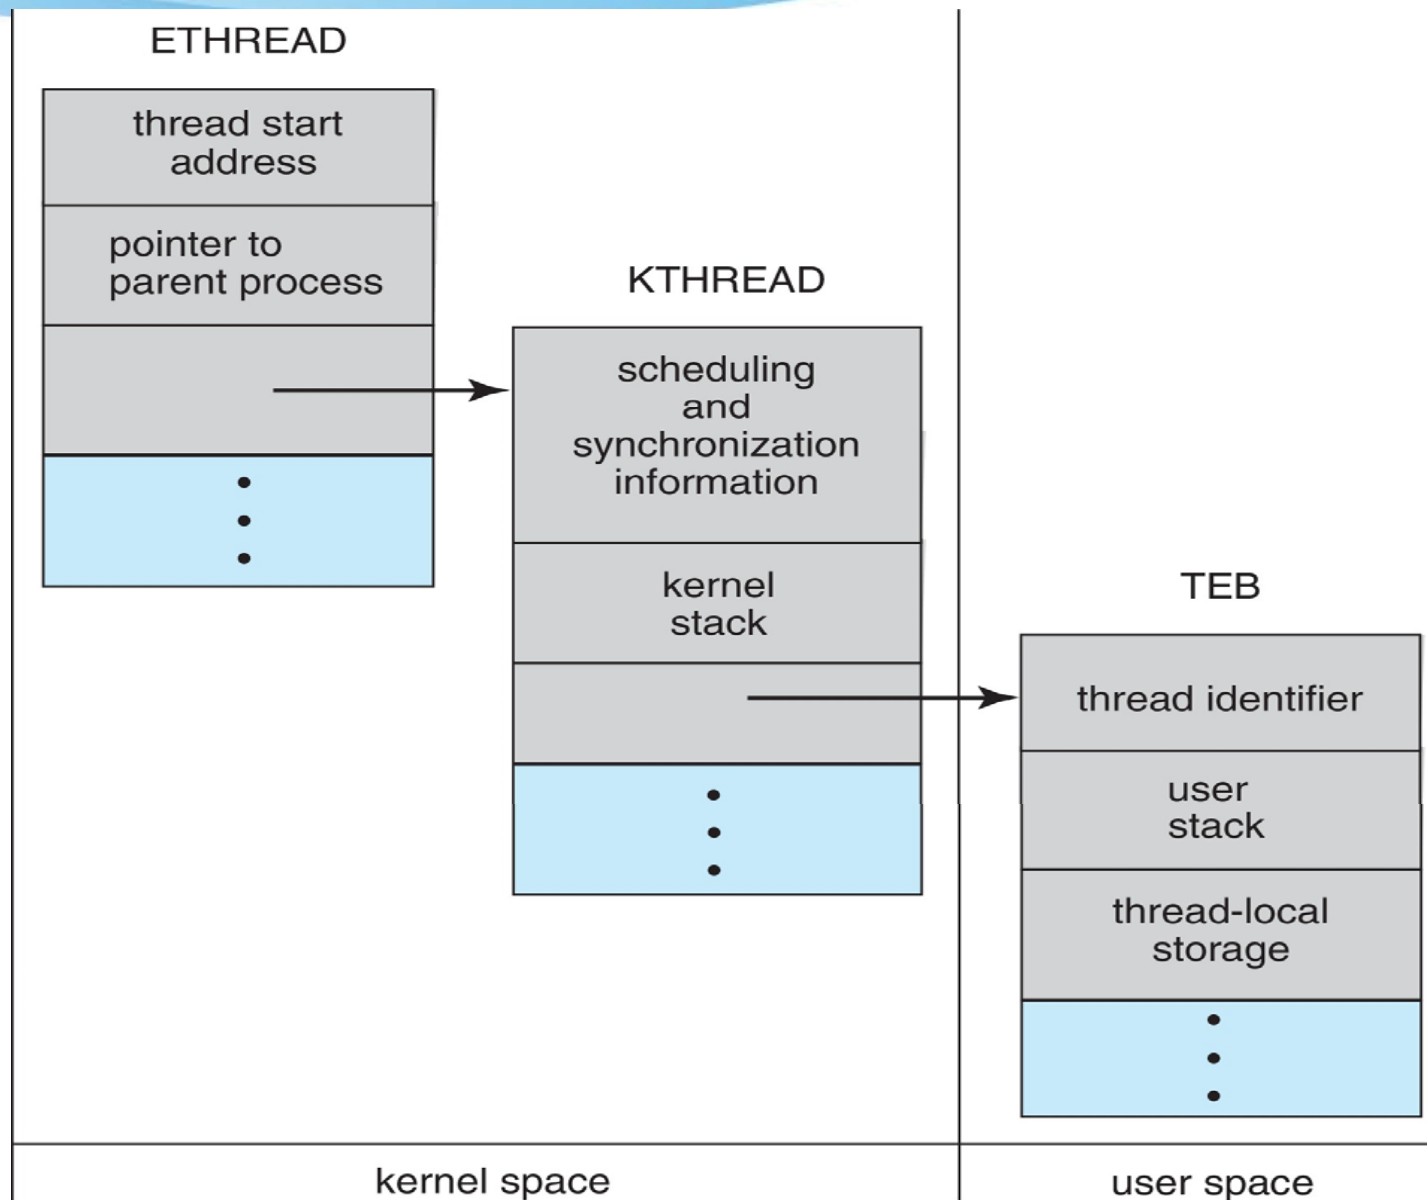
- Explore how threads are implemented in Windows and Linux systems.

Windows Threads



- Implements the one-to-one mapping
 - also support a [fiber library](#) that provides the M:M model
- Each thread contains
 - A thread id
 - Register set
 - Separate user and kernel stacks
 - Private data storage area (i.e. thread specific data)
- The register set, stacks, and private storage area are known as the **context** of the threads
- The primary data structures of a thread include:
 - ETHREAD (executive thread block)
 - KTHREAD (kernel thread block)
 - TEB (thread environment block)

Data Structures of a Windows Thread



Linux Threads

- Linux refers to them as *tasks* rather than *threads* or *processes*
- Thread creation is done through **clone()** system call
- **clone()** allows a child task to share the address space of the parent task (process)

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

Java Threads

- Java threads are managed by the JVM
- Java threads may be created by
 - Extending Thread class
 - Implementing the Runnable interface

Java Thread States

