



Chapter 6

Synchronization Tools

Da-Wei Chang

CSIE.NCKU

*Source: Abraham Silberschatz, Peter B. Galvin, and Greg Gagne,
"Operating System Concepts", 10th Edition, Wiley.*

Outline



- Background
- The Critical-Section Problem
- Peterson's Solution
- Hardware Support for Synchronization
- Mutex Locks
- Semaphores
- Monitors
- Liveness

Background

- **Concurrent** access to **shared** data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the **orderly execution** of cooperating processes
 - **cooperating processes**: processes that access the shared data
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers.
 - use an integer **count**
 - keeps track of the number of full buffers
 - Initially set to 0
 - Incremented by the producer after it produces a new buffer
 - Decrement by the consumer after it consumes a buffer

Producer

```
while (true)
{
    /* produce an item and put in nextProduced */
    while (count == BUFFER_SIZE)
        ; // do nothing
    buffer [in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    count++;
}
```

Consumer

```
while (true)
{
    while (count == 0)
        ; // do nothing
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;
    /* consume the item in nextConsumed */
}
```

Race Condition

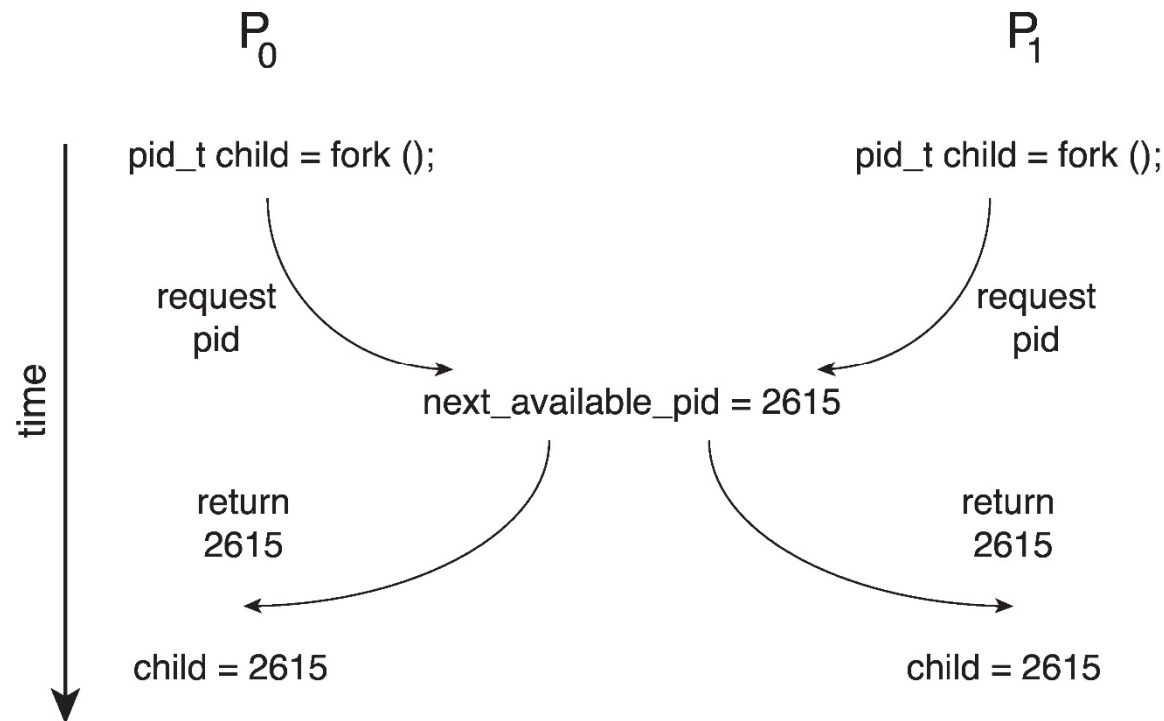
- `count++` could be implemented as

```
register1 = count
register1 = register1 + 1
count = register1
```
 - `count--` could be implemented as

```
register2 = count
register2 = register2 - 1
count = register2
```
 - Consider this execution interleaving with “count = 5” initially:
 - S0: producer execute `register1 = count` {register1 = 5}
 - S1: producer execute `register1 = register1 + 1` {register1 = 6}
 - S2: consumer execute `register2 = count` {register2 = 5}
 - S3: consumer execute `register2 = register2 - 1` {register2 = 4}
 - S4: producer execute `count = register1` {count = 6}
 - S5: consumer execute `count = register2` {count = 4}
- ➔ ***Race condition...***

Race Condition – Another Example

- Processes P_0 and P_1 are creating child processes using the `fork()` system call
- Race condition on kernel variable **next_available_pid** which represents the next available process identifier (pid)



Critical Section

- Consider n cooperating processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has a **critical section**
 - A segment of code, in which the process may update shared data
- *When a process is in the critical section, the others cannot enter their critical sections!*
 - no two processes in the same critical section at the same time
- Critical section problem
 - Design a protocol that processes can use to coordinate
 - Entry section: code to try/request to enter the critical section
 - Exit section: code that follows the critical section

Critical Section

do

{

entry section

critical section *//update shared data here...*

exit section

remainder section

} while (TRUE)

Solution to Critical-Section Problem

A solution should satisfy the following requirements

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the **selection** of the processes that will enter the critical section next **cannot be postponed indefinitely**
3. **Bounded Waiting** - A **bound** must exist on the **number of times** that other processes are allowed to enter their critical sections **after** a process has made a request to enter its critical section and **before** that request is granted – guarantee a given process will finally be selected
 - Assume that each process executes at a **nonzero speed**
 - No assumption concerning relative speed of the N processes

Peterson's Solution

- Two-process solution; a software based solution
- Assume that the LOAD and STORE instructions are atomic; that is, cannot be interrupted.
- The two processes share two variables:
 - int **turn**;
 - Boolean **flag[2]**;
- The variable **turn** indicates whose turn it is to enter the critical section
- The **flag** array is used to indicate whether or not a process is **ready** to enter the critical section
 - **flag[i]** = TRUE implies that process P_i is ready!

Peterson's Solution

Algorithm for Process P_i

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while ( flag[j] && turn == j); // you first
```

CRITICAL SECTION

```
    flag[i] = FALSE; // exit
```

REMAINDER SECTION

```
} while (TRUE);
```

Algorithm for Process P_j

```
do {  
    flag[j] = TRUE;  
    turn = i;  
    while ( flag[i] && turn == i); // you first
```

CRITICAL SECTION

```
    flag[j] = FALSE; // exit
```

REMAINDER SECTION

```
} while (TRUE);
```

Peterson's Solution

- Mutual exclusion is met
 - **turn** can either be **i** or **j**
- Progress and bounded-waiting are met
 - Once P_j exits, it sets $\text{flag}[j]$ to false, allows P_i to enter
 - If P_j resets $\text{flag}[j]$ to true, it must also set turn to i
 - Allow P_i to enter
 - P_i will enter critical section after at most one entry by P_j (bounded-waiting)
 - Selection will not be postponed forever (progress)

Peterson's Solution

- Although useful for demonstrating an algorithm, Peterson's Solution is **not guaranteed to work on modern architectures!**
- Understanding why it will not work is also useful for better understanding race conditions.
- To improve performance, **processors and/or compilers** may reorder operations that have **no dependencies**.
- For single-threaded this is ok as the result will always be the same.
- For multithreaded the reordering may produce inconsistent or unexpected results!

Peterson's Solution

- Two threads share the data

```
boolean flag = false;  
int x = 0;
```

- Thread 1 performs

```
while (!flag);  
print x;
```

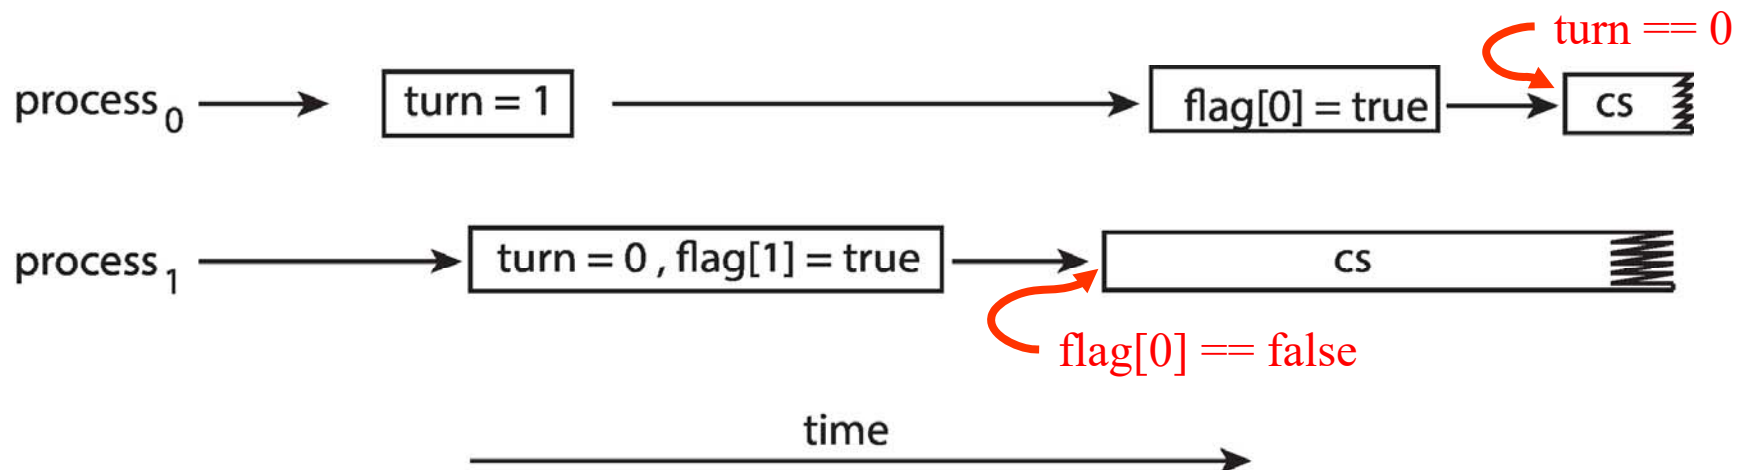
- Thread 2 performs

```
x = 100;  
flag = true;
```

- What is the expected output?
 - 100? right?

Peterson's Solution

- However, the operations for Thread 2 may be reordered:
 `flag = true;`
 `x = 100;`
- If this occurs, the **output may be 0!**
- The effects of instruction reordering in Peterson's Solution



Both processes in their critical sections at the same time!

Hardware Support for Synchronization

- Many systems provide **hardware support** for synchronization
- Uniprocessor – could **disable interrupts**
 - Currently running code would execute without preemption
 - Common in embedded systems
 - However,
 - **Not efficient** on multiprocessor systems
 - Need to ask other processors to disable interrupts
 - Operating systems using this not broadly scalable
 - Clocks may stop in the critical section

Hardware Support for Synchronization



- We will look at three additional forms of hardware support
 1. Memory barriers
 2. Hardware instructions
 3. Atomic variables

Memory Barriers

- Memory models
 - **Strongly ordered:** a memory modification of one processor is immediately visible to all other processors.
 - **Weakly ordered:** a memory modification of one processor may not be immediately visible to all other processors.
- A *memory barrier* (or **memory fence**)
 - an instruction (e.g., mfence for x86) that forces any change in memory to be propagated (made visible) to all other processors.
 - ensure that all loads/stores are **finished** before any subsequent load/store operations are performed

Memory Barrier

- We could add a memory barrier to the following instructions to ensure Thread 1 outputs 100:

- Thread 1 now performs

```
while (!flag)
    memory_barrier(); //flag is loaded before x
print x;
```

- Thread 2 now performs

```
x = 100;
memory_barrier(); // x is assigned before flag
flag = true;
```

Using Memory Barriers for Peterson's Solution

- Place a memory barrier between the first two assignment statements

```
flag[i] = TRUE;
```

```
memory_barrier();
```

```
turn = j;
```

Critical Sections and Locks

```
do
{
    Acquire lock // entry section
    Critical section...
    Release lock // exit section
    Remainder section...
} while (TRUE)
```

Synchronization Hardware



- Modern machines provide special **atomic hardware instructions** for locks
 - Atomic = non-interruptable
- **Test-and-Set** instruction
- **Compare-and-Swap** instruction
- **Swap** instruction

test_and_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
```

```
{  
    boolean rv = *target;  
    *target = true;  
    return rv;  
}
```

Done by a single instruction

- Executed atomically
- Returns the original value of passed parameter
- Set the new value of passed parameter to **true**

Solution using test_and_set

- Shared boolean variable **lock**, initialized to **FALSE**.
- Solution

do {

while (**test_and_set (&lock)**)
; // do nothing

false → true
true → true

..... critical section here

lock = FALSE;

..... remainder section here

} while (TRUE);

compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value)
{
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}
```

Done by a single instruction

- Executed atomically
- Returns the original value of ***value**
- Set the ***value** to the **new_value** if ***value == expected**. That is, the swap takes place only under this condition.

Solution using `compare_and_swap`

- Shared integer **lock**, initialized to 0
- Solution:

```
while (true){  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    ... critical section here...  
  
    lock = 0;  
  
    ... remainder section here ...  
}
```

swap Instruction

- Definition

```
void swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

Done by a single instruction

Solution using swap

- Shared Boolean variable **lock** initialized to **FALSE**; Each process has a **local** Boolean variable **key**.

- Solution:

Used for swapping with *lock*

```
do {  
    key = TRUE;  
    while (key == TRUE)    Swap ( &lock, &key );  
    ..... critical section here  
    lock = FALSE;  
    ..... remainder section here  
} while (TRUE);
```

Lock: TRUE, key: FALSE

Synchronization Hardware

- The above examples
 - do not satisfy bounded waiting...
- The following code satisfy all the three requirements

```
Boolean waiting[n], lock=FALSE;
do {
    waiting[i] = TRUE;                // wish to acquire the lock
    key = TRUE;
    while (waiting[i] && key)  key = TestandSet(&lock); // try to acquire the lock
    waiting[i] = FALSE;
    // .....critical section here.... //
    j = (i+1)%n;
    while ( (j!= i) && !waiting[j] ) j = (j+1)%n; // keep locked and allow the
                                                    // next to come in
    if ( j == i ) lock = FALSE;
    else waiting[j] = FALSE;
    // .....remainder section here.... //
} while (TRUE);
```

Atomic Variables

- Typically, instructions such as compare-and-swap are used as building blocks for other synchronization tools.
- One tool is an **atomic variable** that provides **atomic updates** on **basic data types** such as integers and booleans.
 - Useful when the critical section contains the update of a single shared variable
- For example, the **increment()** operation on the atomic variable k ensures k is incremented without interruption:

increment(& k);

Atomic Variables

- The **increment()** function can be implemented as follows:

```
void increment(atomic_int *v)
{
    int temp;

    do {
        temp = *v;
    }
    while (temp != (compare_and_swap(v,temp,temp+1)));
}
```


***v = 5**

Thread 1	Thread 2
temp = *v = 5	---
---	temp = *v = 5
---	compare_and_swap(v, temp, temp+1); // *v = temp+1 = 6; return 5
compare_and_swap(v, temp, temp+1); // *v != temp, v = 6, return 6;	
temp != 6; while(TRUE), next round	
temp = *v = 6	
	temp == 5; while (FALSE)
	increment() finishes Return to calling procedure
compare_and_swap(v, temp, temp+1); // *v == temp+1=7, return 6;	
temp == 6; while (FALSE)	
increment() finishes Return to calling procedure	

Mutex Locks

- Previous solutions not target for application programmers
- A less complicated syn. interface to programmers - **mutex lock**
- Protect a critical section by first **acquire()** a lock then **release()** the lock
 - **ensure mutual exclusion** (see next slide)
- **acquire()/release()** must be atomic
 - Usually implemented via hardware atomic instructions such as compare-and-swap
- This solution requires **busy waiting**
 - a ***spinlock***

Critical Sections and Locks

```
do
{
    acquire()      // Acquire lock
    Critical section...
    release()     // Release lock
    Remainder section...
} while (TRUE)
```

Mutex Lock Definitions

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

```
release() {  
    available = true;  
}
```

- These two functions must be implemented **atomically**.
- Hardware instructions such as test-and-set/compare-and-swap can be used to implement these functions.

Semaphore

- More general than mutex locks
- Typical semaphore implementations do **NOT** require busy waiting
- Semaphore S – **integer** variable
- Two standard operations modify S : **wait()** and **signal()**
 - Originally called **P()** and **V()**
- Can only be accessed via two indivisible (**atomic**) operations
 - **wait** (S) {
 while $S \leq 0$ // atomic for each iteration
 ; // no-op // iteration: if $S > 0$, $S--$; else, no-op and test again
 $S--$;
}
 - **signal** (S) {
 $S++$; // atomic
}

Semaphore as General Synchronization Tool

- **Counting** semaphore – integer value can range over an unrestricted domain
- **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
 - Same as **mutex locks**
- A counting semaphore **S** can be used to implement a binary semaphore

Semaphore as General Synchronization Tool

- Provides mutual exclusion
 - Semaphore S; // initialized to 1
 - wait (S);
Critical Section
signal (S);
- Other usages on semaphores
 - Control the usage of resources with N instances
 - S is initialized to N
 - Synchronization between processes
 - S is initialized to 0

Semaphore Implementation



- Previous semaphore definition requires busy waiting...
- Note that applications may spend lots of time in critical sections. In this case, busy waiting is not a good approach.
- Solution
 - Sleep/block instead of busy waiting

Semaphore Implementation **WITHOUT** Busy Waiting

- In addition to the semaphore **value**, each semaphore has an associated **waiting queue**
- Require blocking/waking-up a process
 - **block**
 - switch process state to “waiting”
 - remove the process from the ready queue
 - **wakeup**
 - switch process state to “ready”
 - place the process in the ready queue
- Both operations may lead to the invocation of the scheduler

Semaphore Implementation **WITHOUT** Busy waiting (Cont.)

- Implementation of wait:

```
wait (S){  
    value--;  
    if (value < 0) { // the value can be negative → # of processes waiting  
        add this process to the waiting queue  
        block(); /* suspend the process */  
    }  
}
```

- Implementation of signal:

```
signal (S){  
    value++;  
    if (value <= 0) {  
        remove a process P from the waiting queue  
        wakeup(P);  
    }  
}
```

Semaphore Implementation

- The list of waiting processes?
 - A list of PCBs
 - FIFO → ensure bounded waiting
- Must guarantee that no two processes can execute `wait()` and `signal()` on the same semaphore at the same time
 - The implementation in the previous slide should be **atomic**
- Thus, implementation becomes the **critical section problem** where the ***wait and signal code are placed in the critical section.***
 - For uniprocessor → disable interrupt
 - For MP → typically use a **busy waiting** approach., e.g. `test_and_set`
 - Acceptable for busy waiting in critical sections of the `wait()` and `signal()` operations since they are short...

Problems with Semaphores

- Using semaphores incorrectly can result in bugs that are hard to detect
- Incorrect use of semaphore operations:
 - `signal(mutex) wait(mutex)` → race
 - `wait(mutex) ... wait(mutex)` → deadlock (*described later*)
 - Omitting of `wait(mutex)` or `signal(mutex)` (or both) → race or deadlock
- The bugs are hard to detect
 - Faults don't always take place

Monitors


- A high-level abstraction that provides a convenient and effective mechanism for process synchronization

```
monitor monitor-name
{
    ... shared variable declarations

    procedure P1 (...) { .... }
    ...
    procedure Pn (...) {.....}
    Initialization code ( ....) { ... }
    ...
}
```

- Programmers provide the operations
- Monitor makes sure **only one process** may be **active within the monitor at a time**

Schematic View of a Monitor



A list of processes waiting for entering the monitor

Monitors


- With monitors, programmers do not need to code the synchronization constraint
- However, the function provided by monitors is limited
 - In many cases, programmers still need to define additional syn. mechanisms
 - Use condition variables

Condition Variables

- condition `x, y`;
- Two operations on a condition variable
 - `x.wait()`
 - a process that invokes the operation is suspended
 - `x.signal()`
 - resumes **one** of processes (if any) that invoked `x.wait ()`
 - No effect if no one is waiting
- No values, just **suspend** and **resume**
 - Different from semaphores

Using them incorrectly are still hard-to-be-detected bugs

Monitor with Condition Variables



If Q wait X and then P signal X

- P is now in monitor, but it will make Q become active in the monitor
- Two possibilities
 - A. **signal & wait** : Q becomes active immediately, P leaves
 - B. **signal & continue**: Q becomes active after P leaves

Implementing a Monitor Using Semaphores

- Use a semaphore **mutex** (init. as 1) for each monitor
 - wait(mutex) before entering the monitor
 - signal(mutex) after leaving the monitor
- A signaling process may wait until the resumed process to leave or wait
 - Use another semaphore **next** (init. as 0)
 - Signaling process can suspend on it
 - **Next_count**
 - Count the number of processes suspend on **next**

Implementing a Monitor Using Semaphores

- Each external procedure F is replaced by

wait (mutex);

F

if (next_count > 0)

signal(next);



Wakeup a signaling process

else

signal(mutex);



Allow another process to
come in

- Each condition variable is implemented by using
 - x_sem : semaphore corresponds to x (init. as 0)
 - x_count: number of processes waiting on x_sem

Implementing a Monitor Using Semaphores

x.wait()

{

x_count++;

if (next_count > 0)

else

wait(x_sem);

x_count--;

}

signal(next);

signal(mutex);

← wait here....

← have been waken

x.signal()

{

if (x_count > 0)

 {

next_count++;

signal(x_sem);

wait(next);

next_count--;

 }

}

← wait here....

← have been waken

Process Resuming Order



- Which process (waiting for the monitor) should be resumed first?
 - FCFS
 - Priority
 - Users can specify priority on suspension
 - `x.wait(priority)`
 - Process with the highest priority is scheduled next

Single Resource Allocation

- Allocate a single resource among competing processes using priority numbers that specify the maximum time a process plans to use the resource

R.acquire(t);

...

access the resource;

...

R.release();

- Where R is an instance of type **ResourceAllocator**
 - *See next slide*

A Monitor to Allocate Single Resource

monitor ResourceAllocator

```
{  
    boolean busy; //whether resource is busy or not  
    condition x;  // wait on x if resource busy  
  
    void acquire (int time) {  
        if (busy)    x.wait(time);  
        busy = true;  
    }  
  
    void release () {  
        busy = FALSE;  
        x.signal();  
    }  
    initialization code() { busy = false; }  
}
```

Incorrect Use of Monitor

- Access a resource without being granted
- Never release the resource
- Release a resource that the process doesn't have
 - Release the resource twice
- Monitors make sure only one process is in it at a time, but it can not solve all the problems related to resource allocation/deallocation

Liveness

- Processes may have to wait indefinitely while trying to acquire a lock
- Waiting indefinitely violates the *progress* and *bounded-waiting* criteria discussed at the beginning of this chapter.
- **Liveness** refers to a set of properties that a system must satisfy to ensure processes make progress.
- Examples of liveness failure
 - Deadlock
 - Starvation
 - Priority inversion

Liveness Failures

- **Deadlock** – two or more processes are waiting infinitely for an event that can be caused by only one of the waiting processes
- Let S and Q be two semaphores initialized to 1

P_0	P_1
wait (S);	wait (Q);
wait (Q);	wait (S);
.	.
.	.
.	.
signal (S);	signal (Q);
signal (Q);	signal (S);

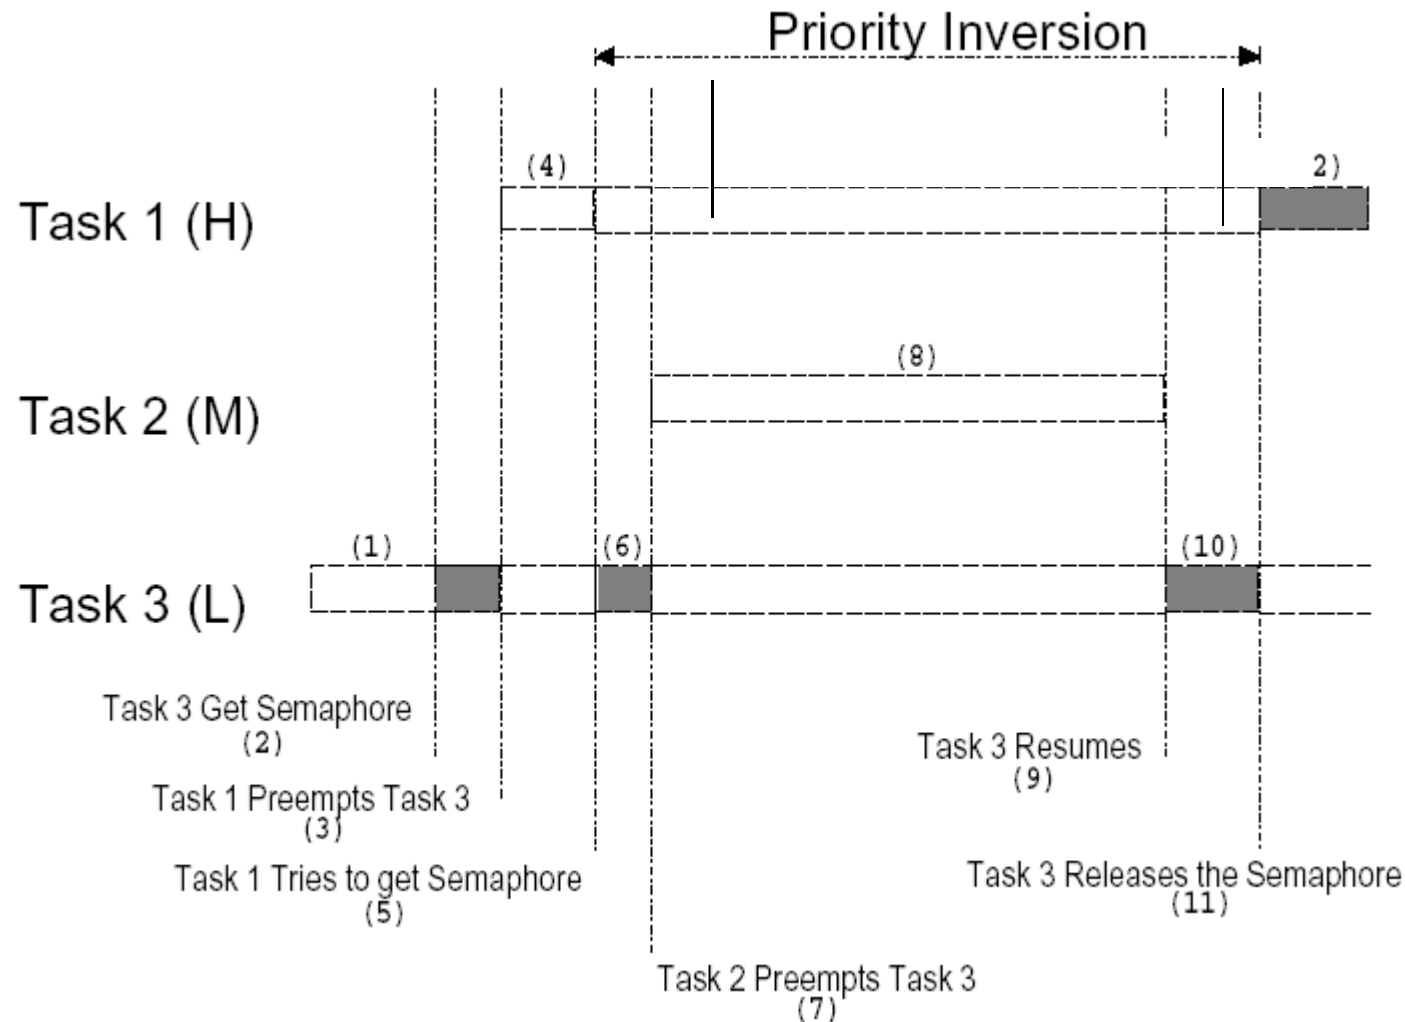
- Consider if P_0 executes wait(S) and P_1 wait(Q). When P_0 executes wait(Q), it must wait for P_1 . However, P_1 also waits for P_0 when it executes wait(S).

Liveness Failures


- **Starvation** – infinite blocking. A process may never be removed from the semaphore queue in which it is suspended.
- **Priority Inversion** – scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - can be solved by **priority-inheritance protocol**

Priority Inversion

- A problem in real-time systems



Priority Inheritance Protocol (PIP)



Some level of priority inversion can not be avoided!!