




成功大學

National Cheng Kung University


Chapter 4

Pipelined Processor




Outline

- A pipelined datapath
- Pipelined control
- Data hazards and forwarding
- Data hazards and stalls
- Branch hazards




Pipelining Overview

- Assume four steps to do laundry, each one take 30min
 - Place the clothes in the washer
 - Place the clothes in the dryer
 - Fold the clothes
 - Ask roommate to put the clothes away
- Nonpipelined laundry: one at a time




- Time for four load
 - $2 \text{ hr} * 4 = 8 \text{ hr}$

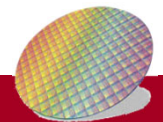


Pipelining Analogy

- **Pipelined** laundry: overlap the task
 - Parallelism improves performance




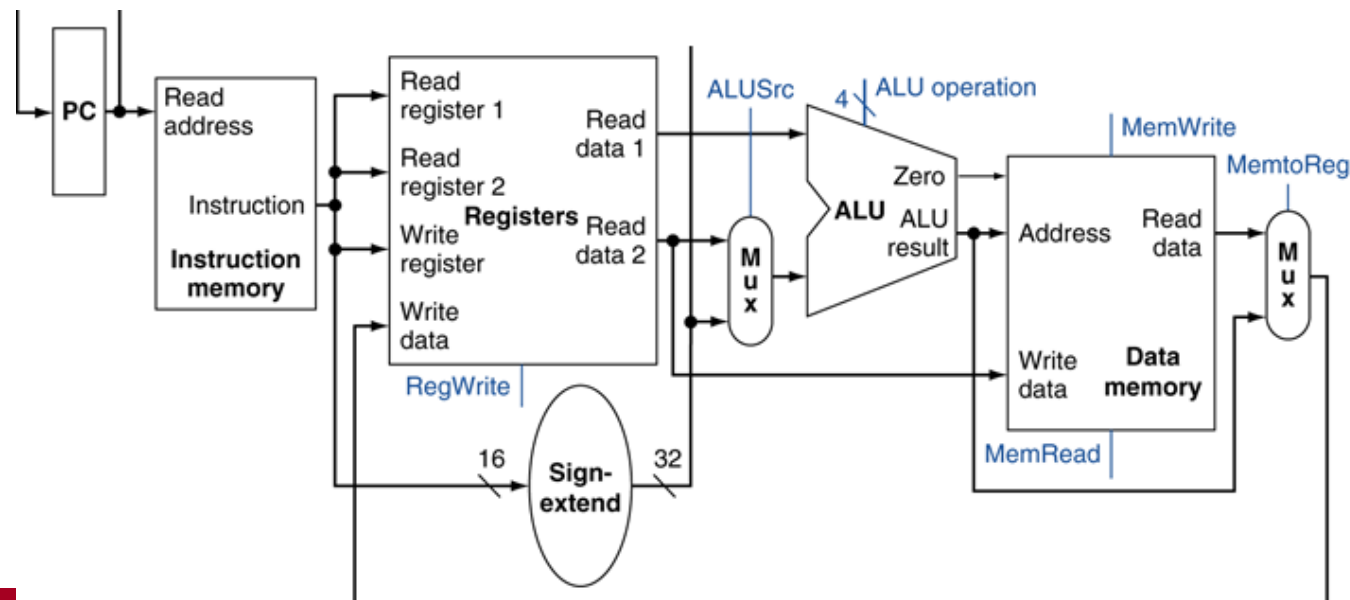
- Time for four loads
 - 3.5 hr
- Speedup
 - $= 8 / 3.5 = 2.3$
- Non-stop for a lot of loads
 - Speedup
 - $= 2n / 0.5n + 1.5 \approx 4$
 - = number of stages





MIPS Pipeline

- Five stages, one step per stage
 1. **IF**: Instruction fetch from memory
 2. **ID**: Instruction decode & register read
 3. **EX**: Execute operation or calculate address
 4. **MEM**: Access memory operand (for **lw** and **sw** instruction)
 5. **WB**: Write result back to register




Pipeline Performance

- Assume time for stages is
 - 100ps for register read or write
 - 200ps for other stages
- Time comparison between pipelined with single-cycle datapath

Longest data path:
800ps

Required time for single-cycle datapath


Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps



Pipeline Performance


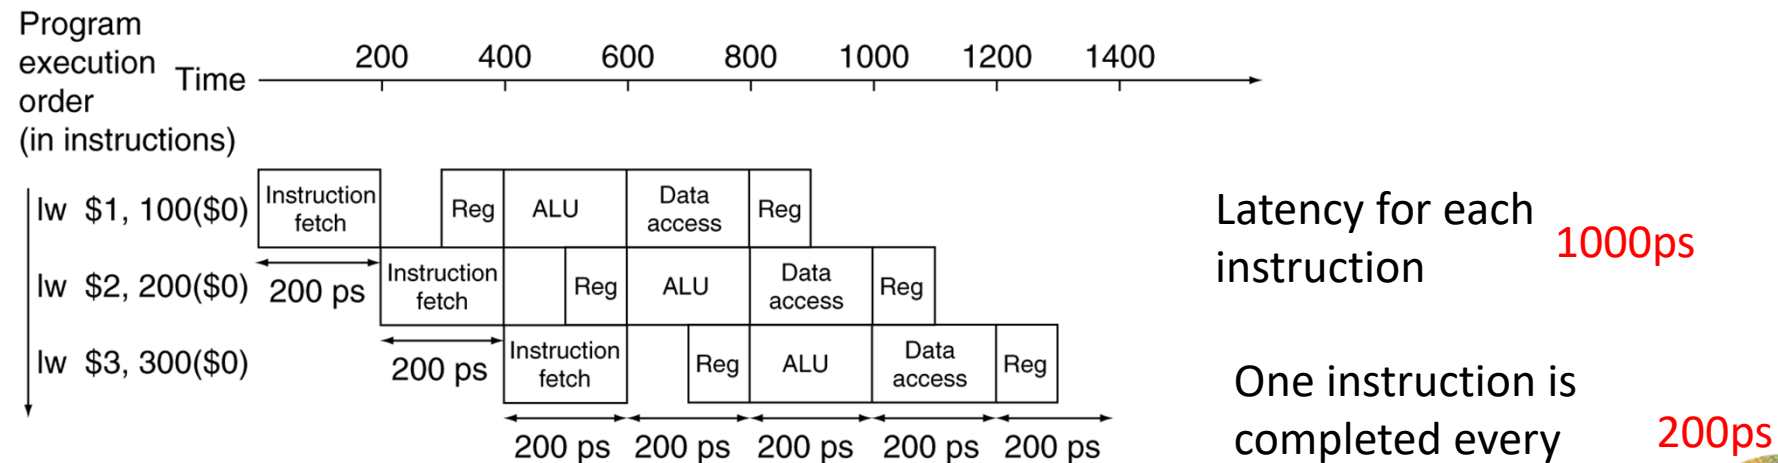
Single-cycle

($T_c = 800\text{ps}$)



Pipelined

($T_c = 200\text{ps}$)




Pipeline Speedup

- Stage may be **unbalanced**: may waste time in a stage
- If all stages are balanced (all take the same time)


$$\frac{\text{Time Between instructions}_{\text{pipelined}}}{\text{Time between instruction}_{\text{nonpipelined}}} = \frac{1000/5}{\text{Number of Stages}} = 200$$

- Speedup is due to increased throughput
 - **Latency** (time for each instruction) does **not decrease**
 - Latency is 1000 ps, but one instruction is completed every 200ps



Pipeline Hazards

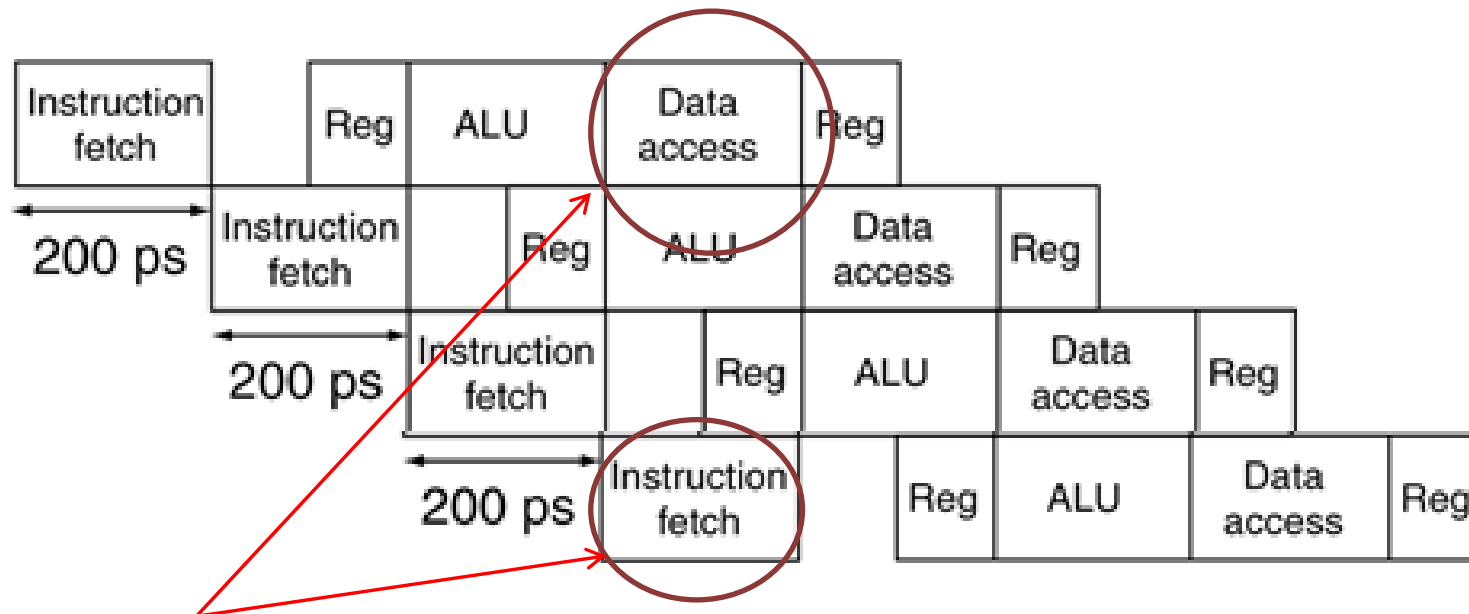
- Hazards are **situations** that prevent **starting** the next instruction in the next cycle
- **Structure** hazards
 - A required **resource** is **busy**. Need to **wait** until the resource is ready
 - E.g. instruction and data memory are separated to avoid structure hazards (see next slide)
- **Data** hazard
 - Need to wait for **previous instruction to complete its data** read/write
- **Control** hazard
 - Deciding on control **action** depends on **previous** instruction






Structure Hazards

- Conflict for use of a resource
- Assume only a single memory in MIPS pipeline
 - Load/store and Instruction fetch both require memory access
 - *Stall* will occur if both happen in the **same cycle**
 - **Structure Hazards**


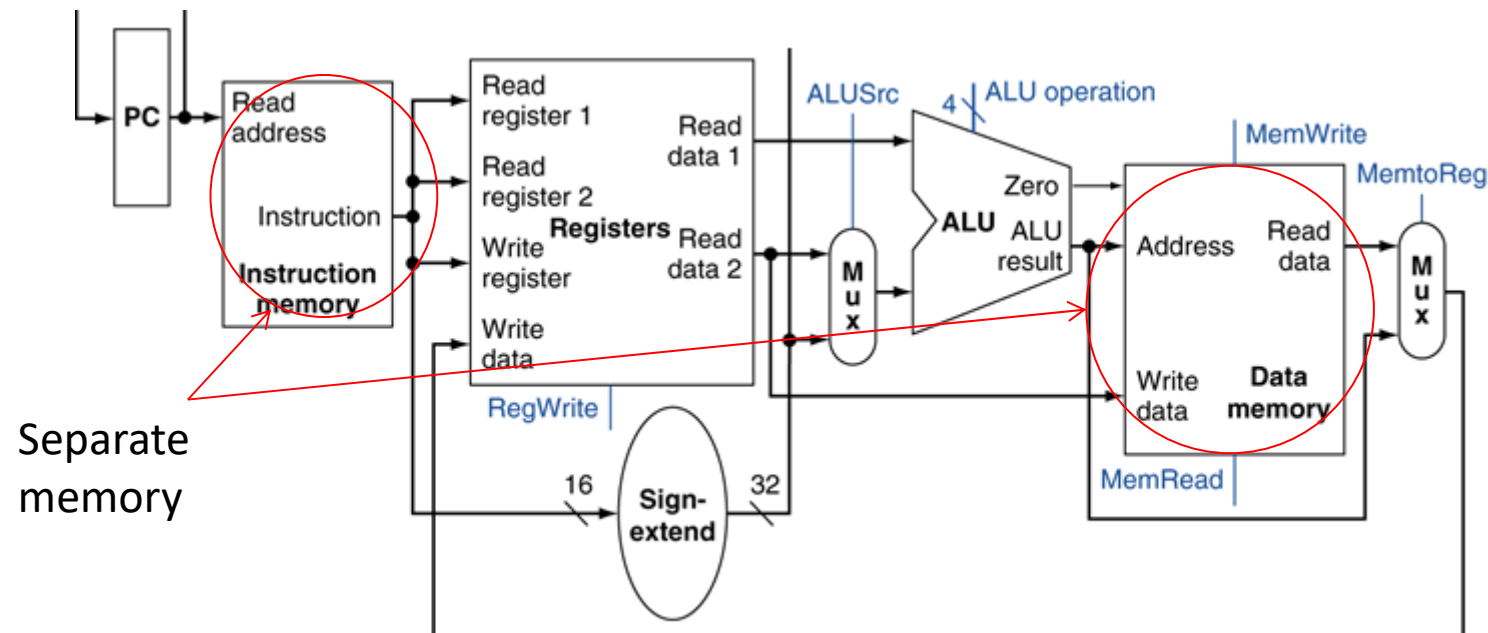


Conflict , both need to access memory at the same time





Solution for Structure Hazards

- Hence, pipelined datapath requires **separate instruction/data** memories
 - Or separate instruction/data caches



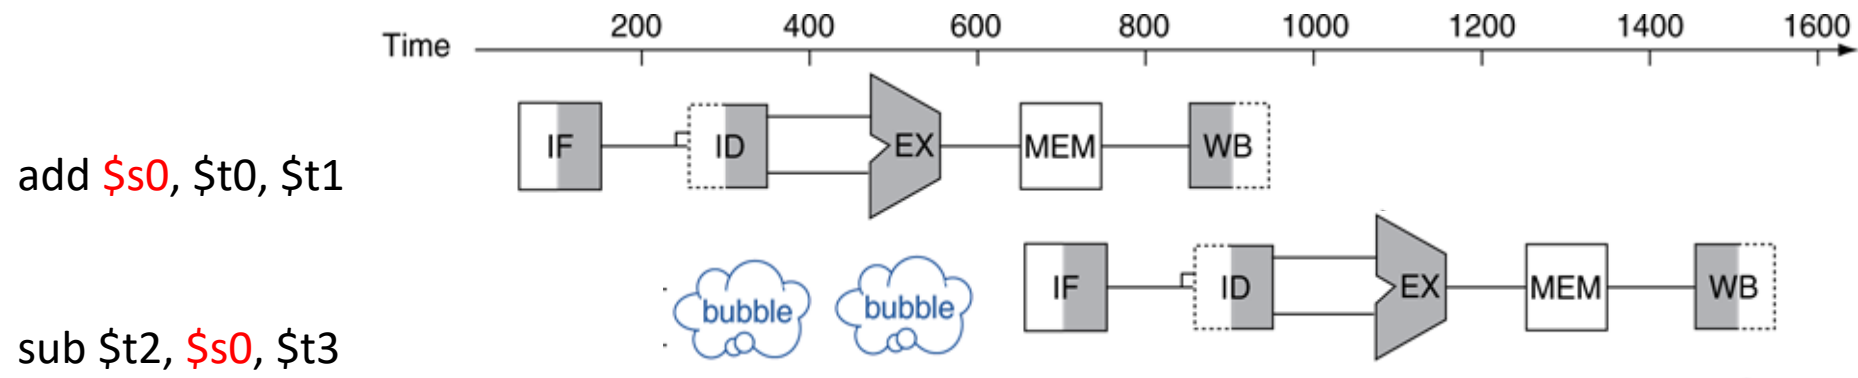
Data Hazards

- An instruction depends on completion of data access by a previous instruction
 - add **\$s0**, \$t0, \$t1
 - sub \$t2, **\$s0**, \$t3




Data Hazards causes stall

- Instruction stall for cycles – waste cycle time
 - Also called pipeline “bubble”




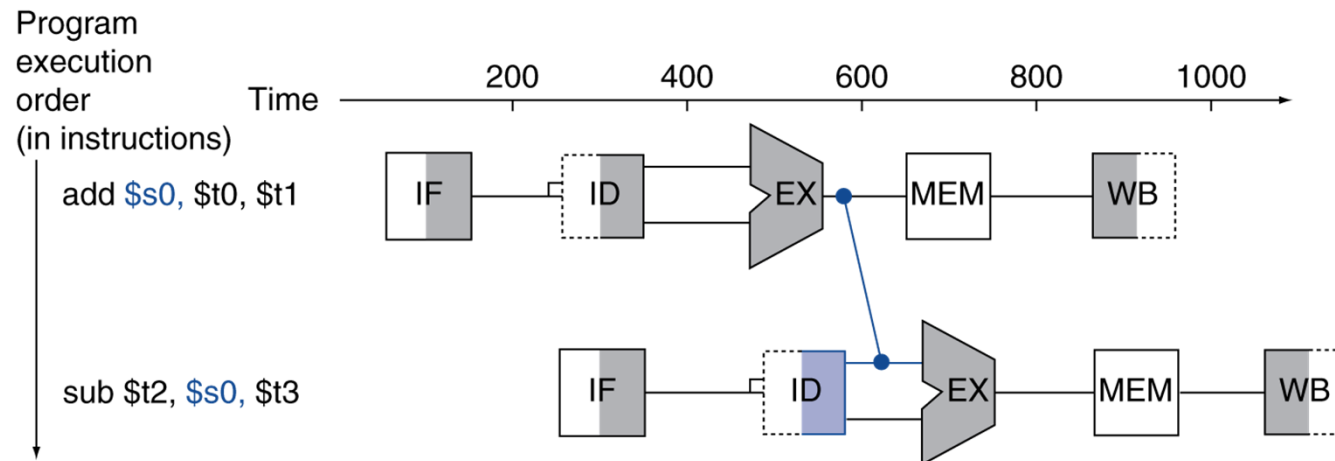
2 bubbles(nop) are added

Can we have a better solution?




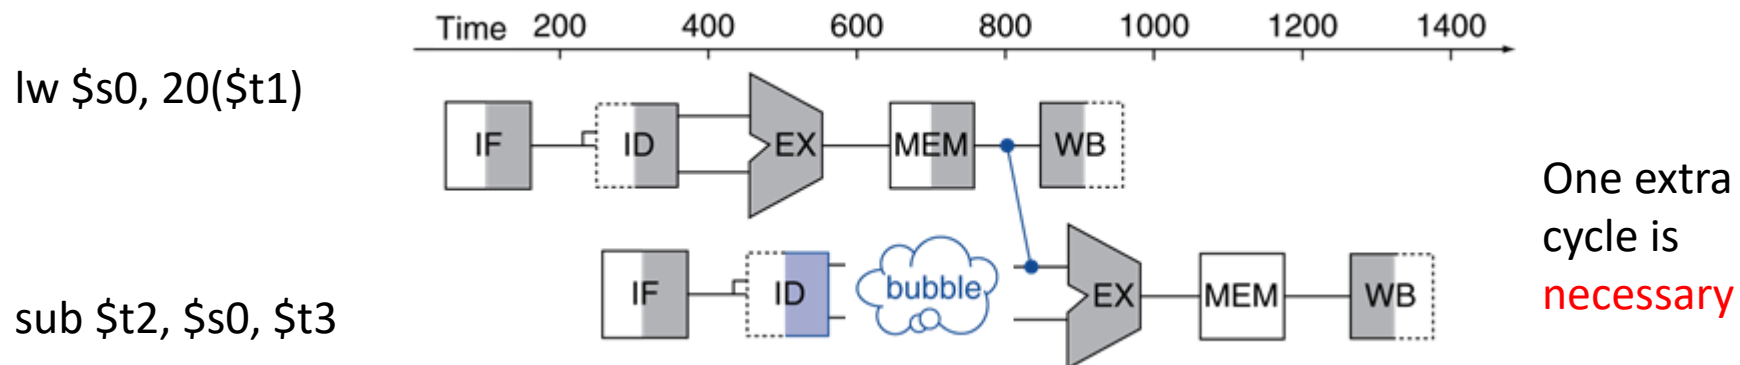
Data Hazard Solution 1 - Forwarding

- Use result when it is **computed**
 - Don't wait for it to be stored in a register
 - Requires extra **connections** in the datapath
 - Also call **bypassing**



Load-Use Data Hazard


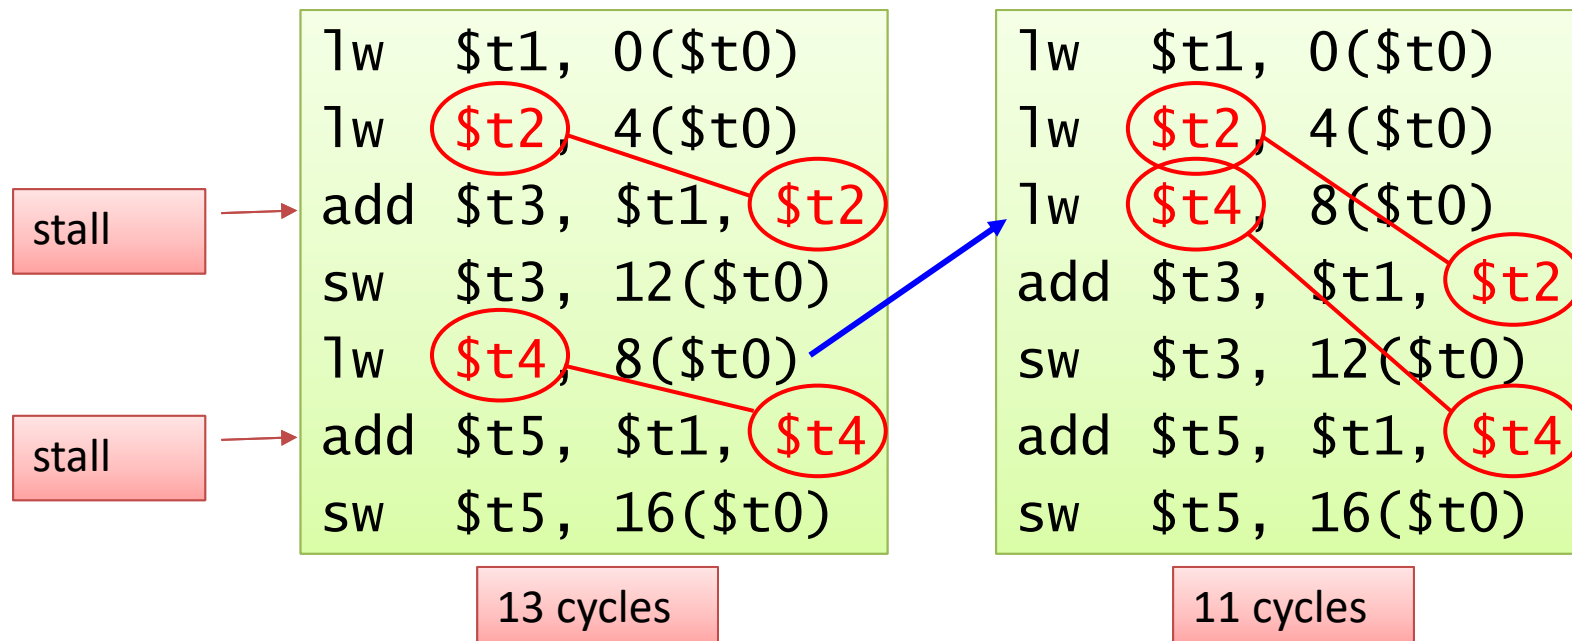
- Forwarding can't always avoid stalls
 - If value is not **ready** when needed
 - Can't forward backward in time!
 - For example, the **lw** instruction, data is ready after **memory** access





Data hazard solution 2: Code Scheduling


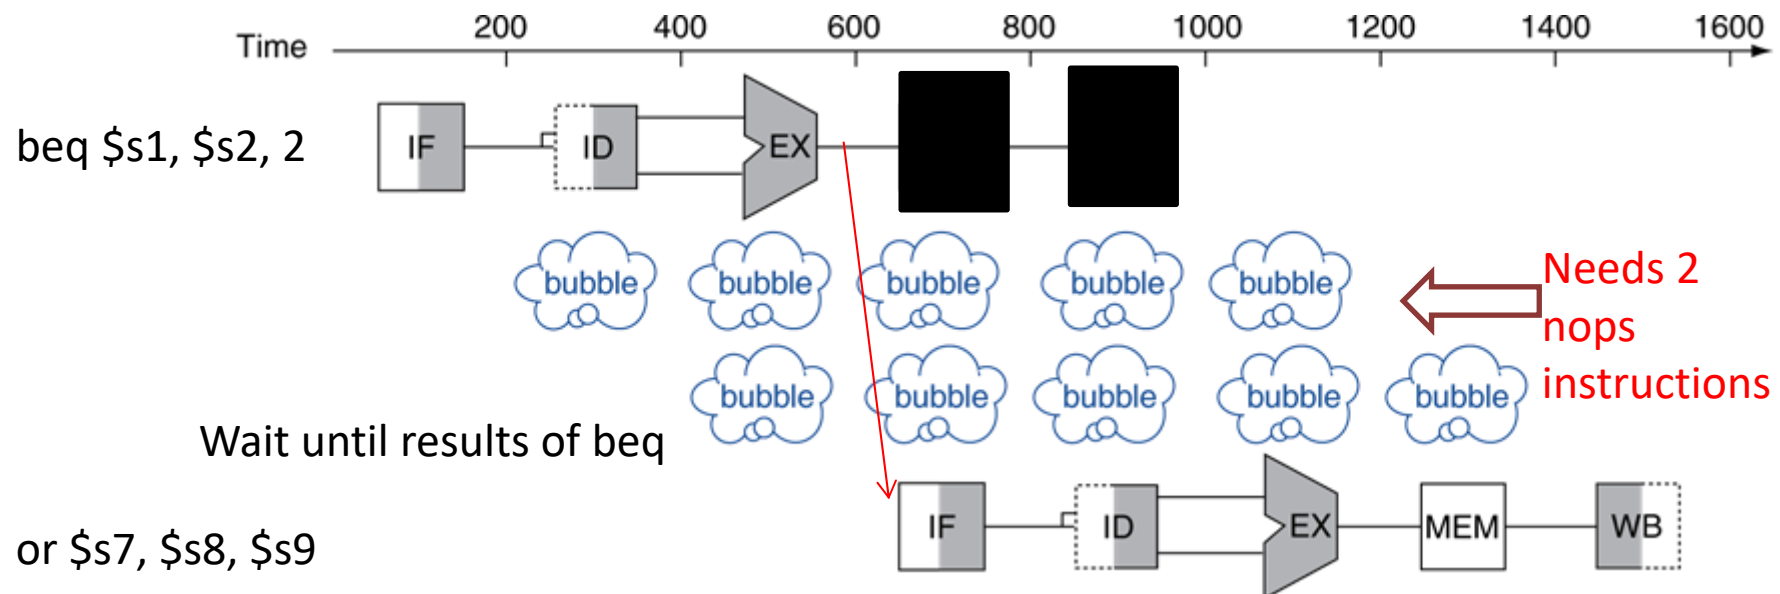
- **Reorder** code to avoid use of load result in the next instruction
- C code for $A = B + E$; $C = B + F$;





Control Hazard causes stalls (bubble)

- **Branch** determines **flow** of control
 - Fetching next **instruction** depends on branch **outcome**
 - Result of Beq is known at the **EX** stage
- Wait until **branch** outcome is determined before fetching next instruction


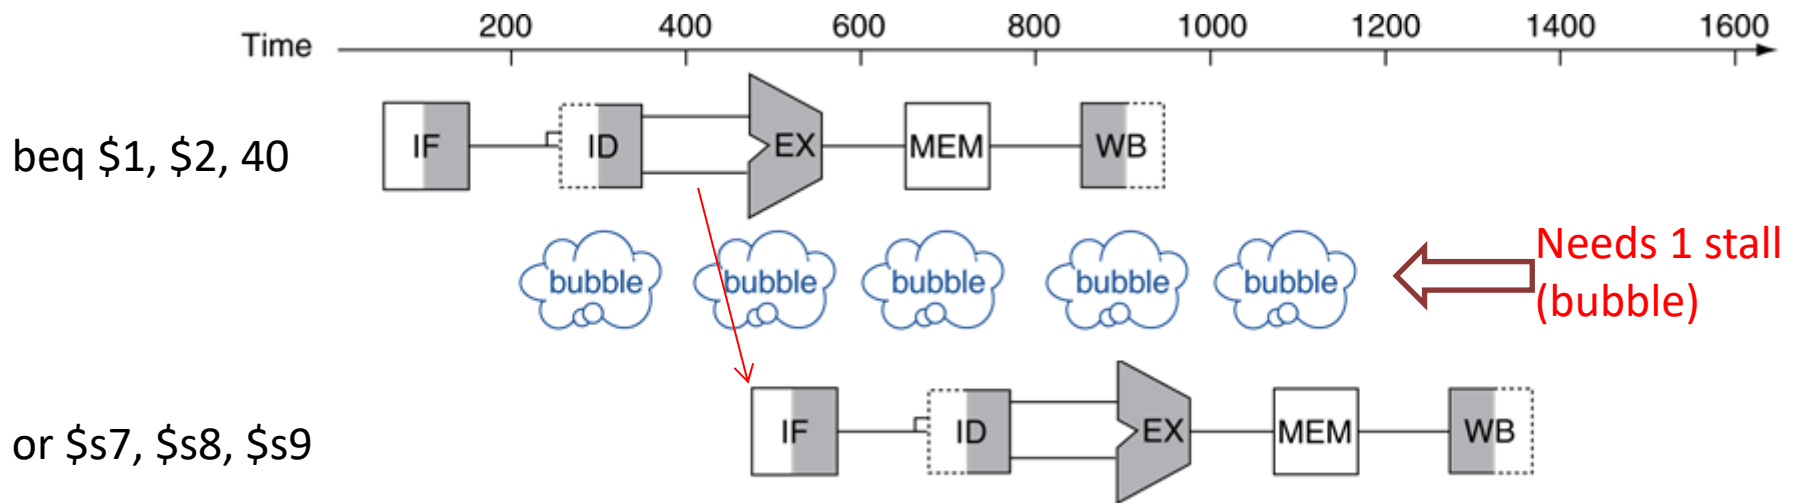


Control Hazard Solution 1: Add hardware in ID stage



成功大學

- Compare **registers** and compute target **early** in the pipeline (Details are in § 4.8)
- **Extra hardware** to **test** registers, calculate the branch address, and update the PC during the **ID stage** of the pipeline => stall is reduced to only **one**







Control Hazard Solution 2: Branch Prediction

- Predict outcome of branch
 - Correct prediction, **no** bubble is added
 - Incorrect prediction, clean incorrect and add **bubble**
- Example: predict branches **not taken**
 - **Fetch** instruction after branch, with no delay
 - If prediction is **correct**, **no** stall
 - If prediction is **incorrect**, ignore the incorrect result, and load **correct** instruction

Program execution order



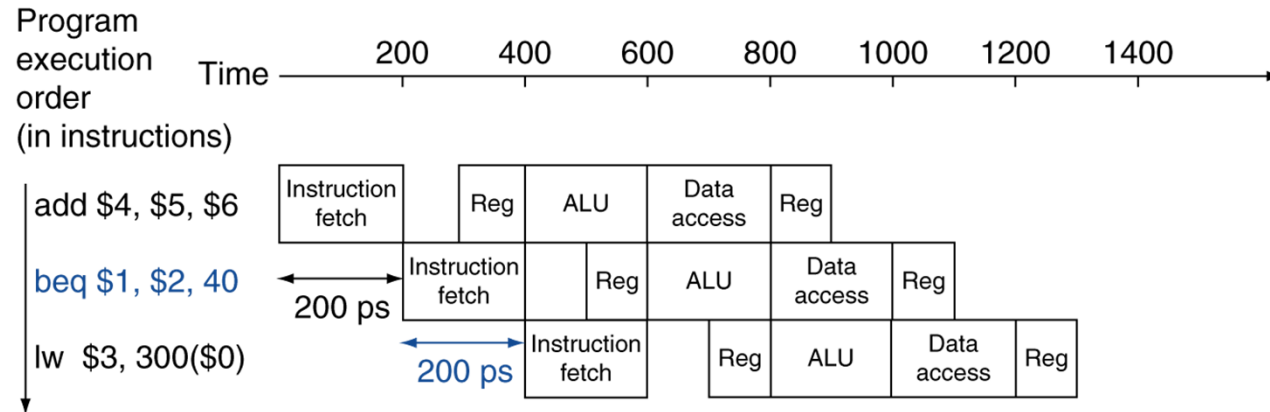
See next **slide** for figures, and more details in § 4.8



MIPS with Predict Not Taken


Branch is determined at ID stage

Correct Prediction




No bubble

Incorrect Prediction



one bubble

Branch is determined



Improve Branch Prediction Accuracy



- **Static** branch prediction
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - **backward** branches => predict **taken**
 - **forward** branches => **not** taken

Backward => predicted Taken

- **Dynamic** branch prediction
 - Hardware measures **actual** branch behaviors
 - Prediction based on **recent histories** of each branch

e.g. Taken, Taken, Taken, Taken

What is the next prediction? Taken for Not Taken ?







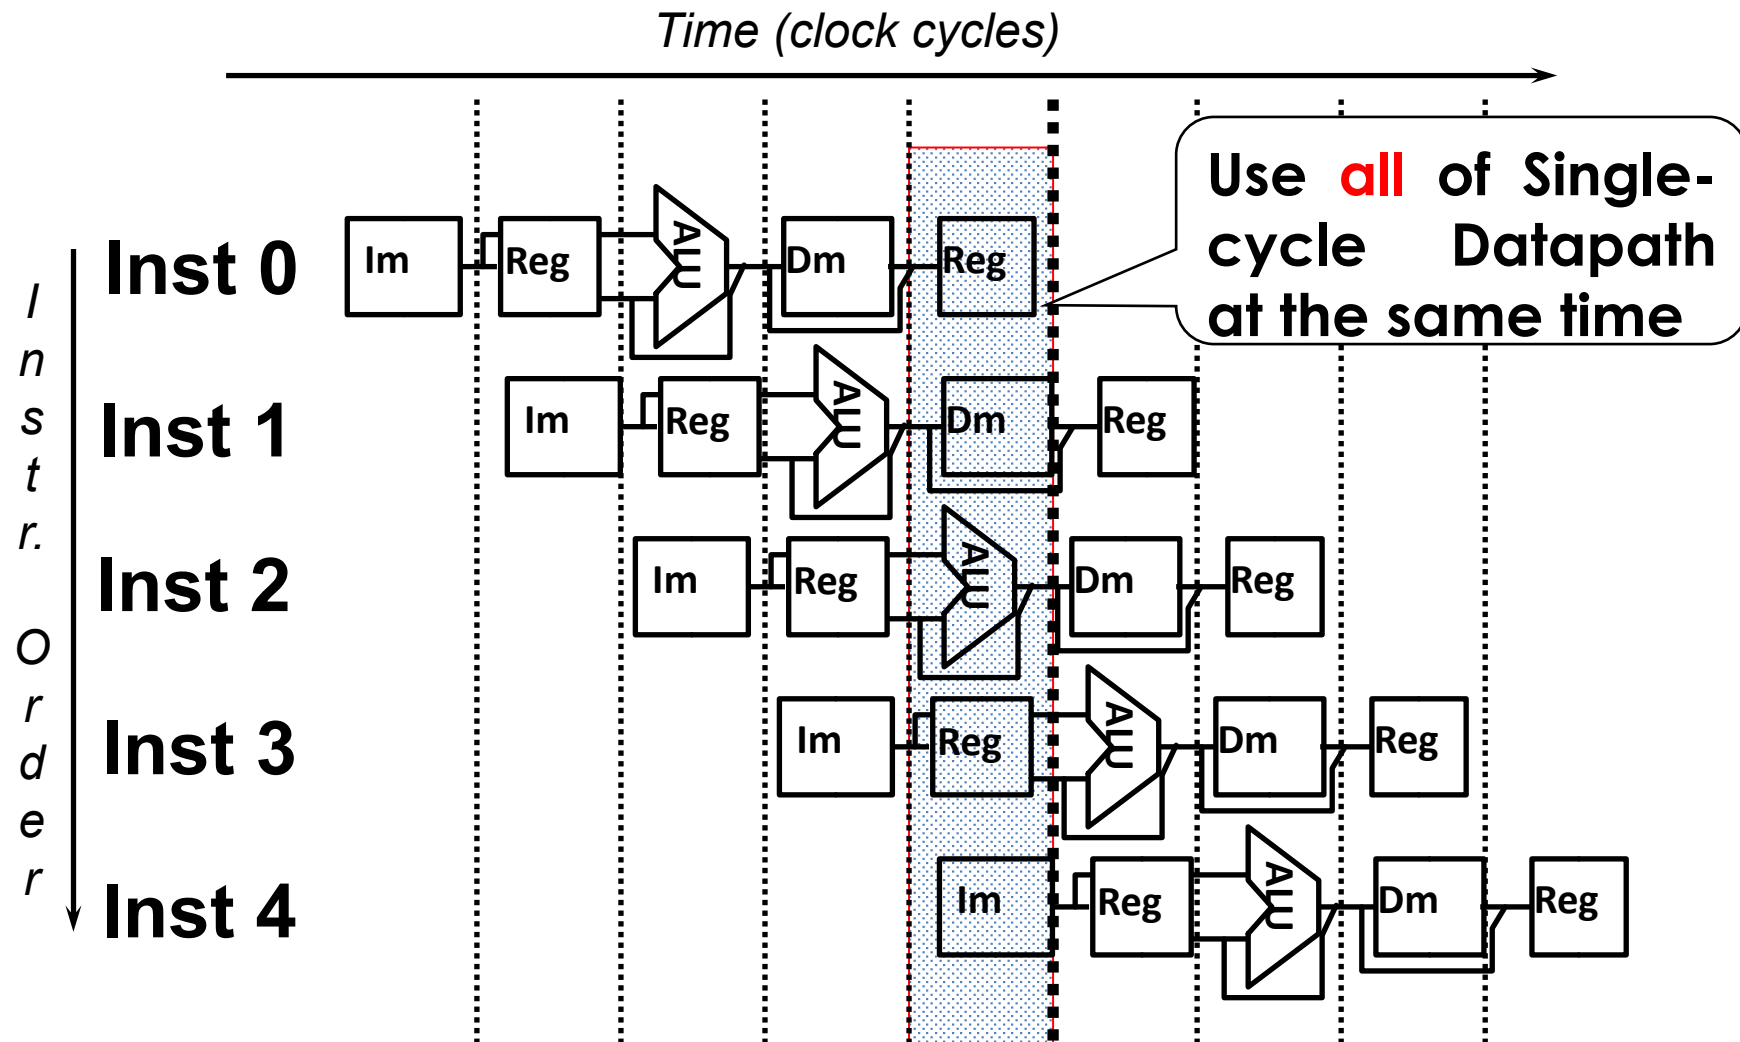
成功大學

National Cheng Kung University

§4.6 Pipelined Datapath and Control




Why Pipeline? Because the Resources are There!


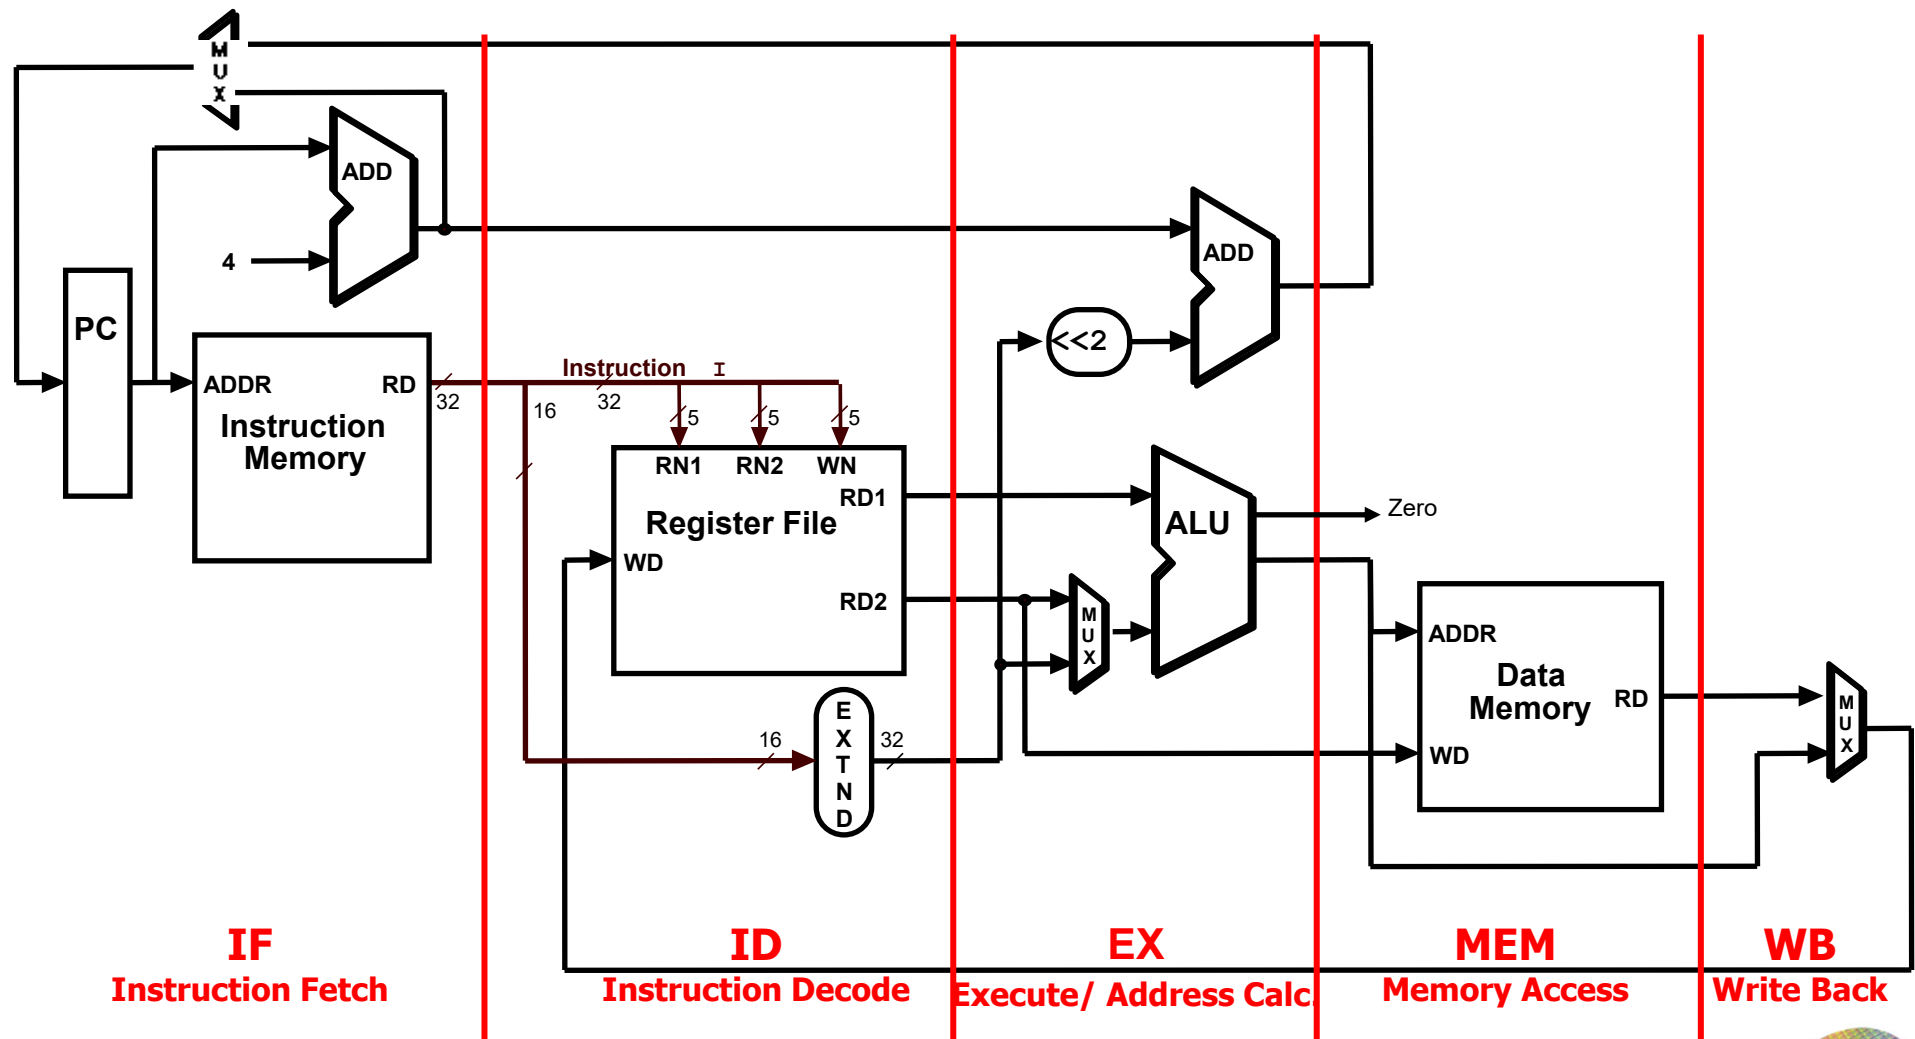


Recall: 5 steps in Instruction Execution

1. Instruction Fetch & PC Increment (IF)
2. Instruction Decode and Register Read (ID)
3. Execution or calculate address (EX)
4. Memory access (MEM)
5. Write result into register (WB)

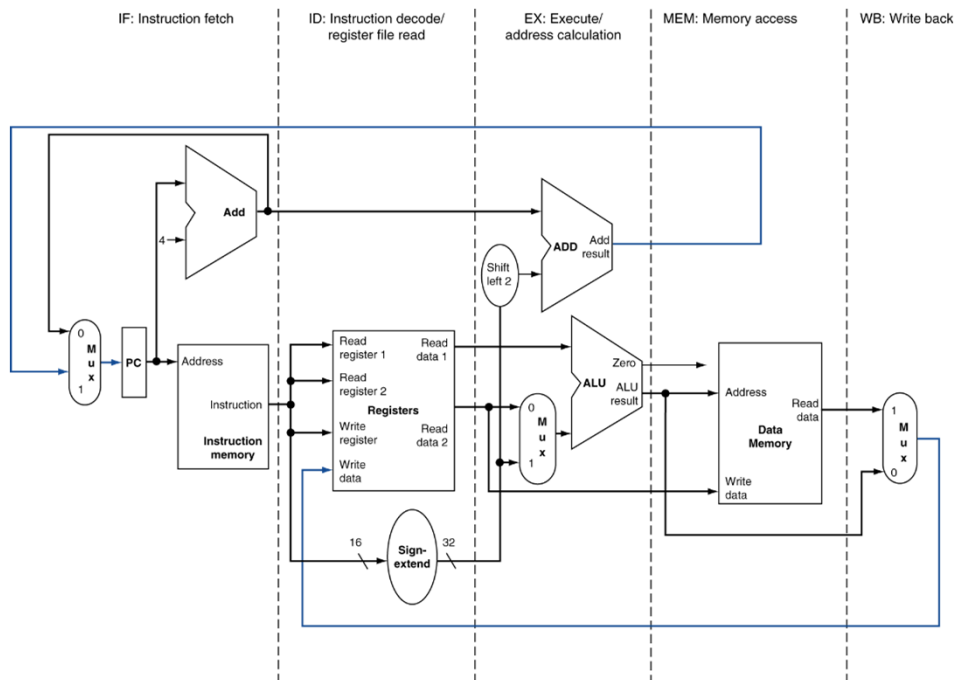


Recall “Single-Cycle” Datapath




Pipelined Datapath

- Break instructions into the following *potential* execution steps
 1. Instruction fetch and PC increment (IF)
 2. Instruction decode and register fetch (ID)
 3. Execution, memory address computation, or branch completion (EX)
 4. Memory access or R-type instruction completion (MEM)
 5. Memory read completion (WB)



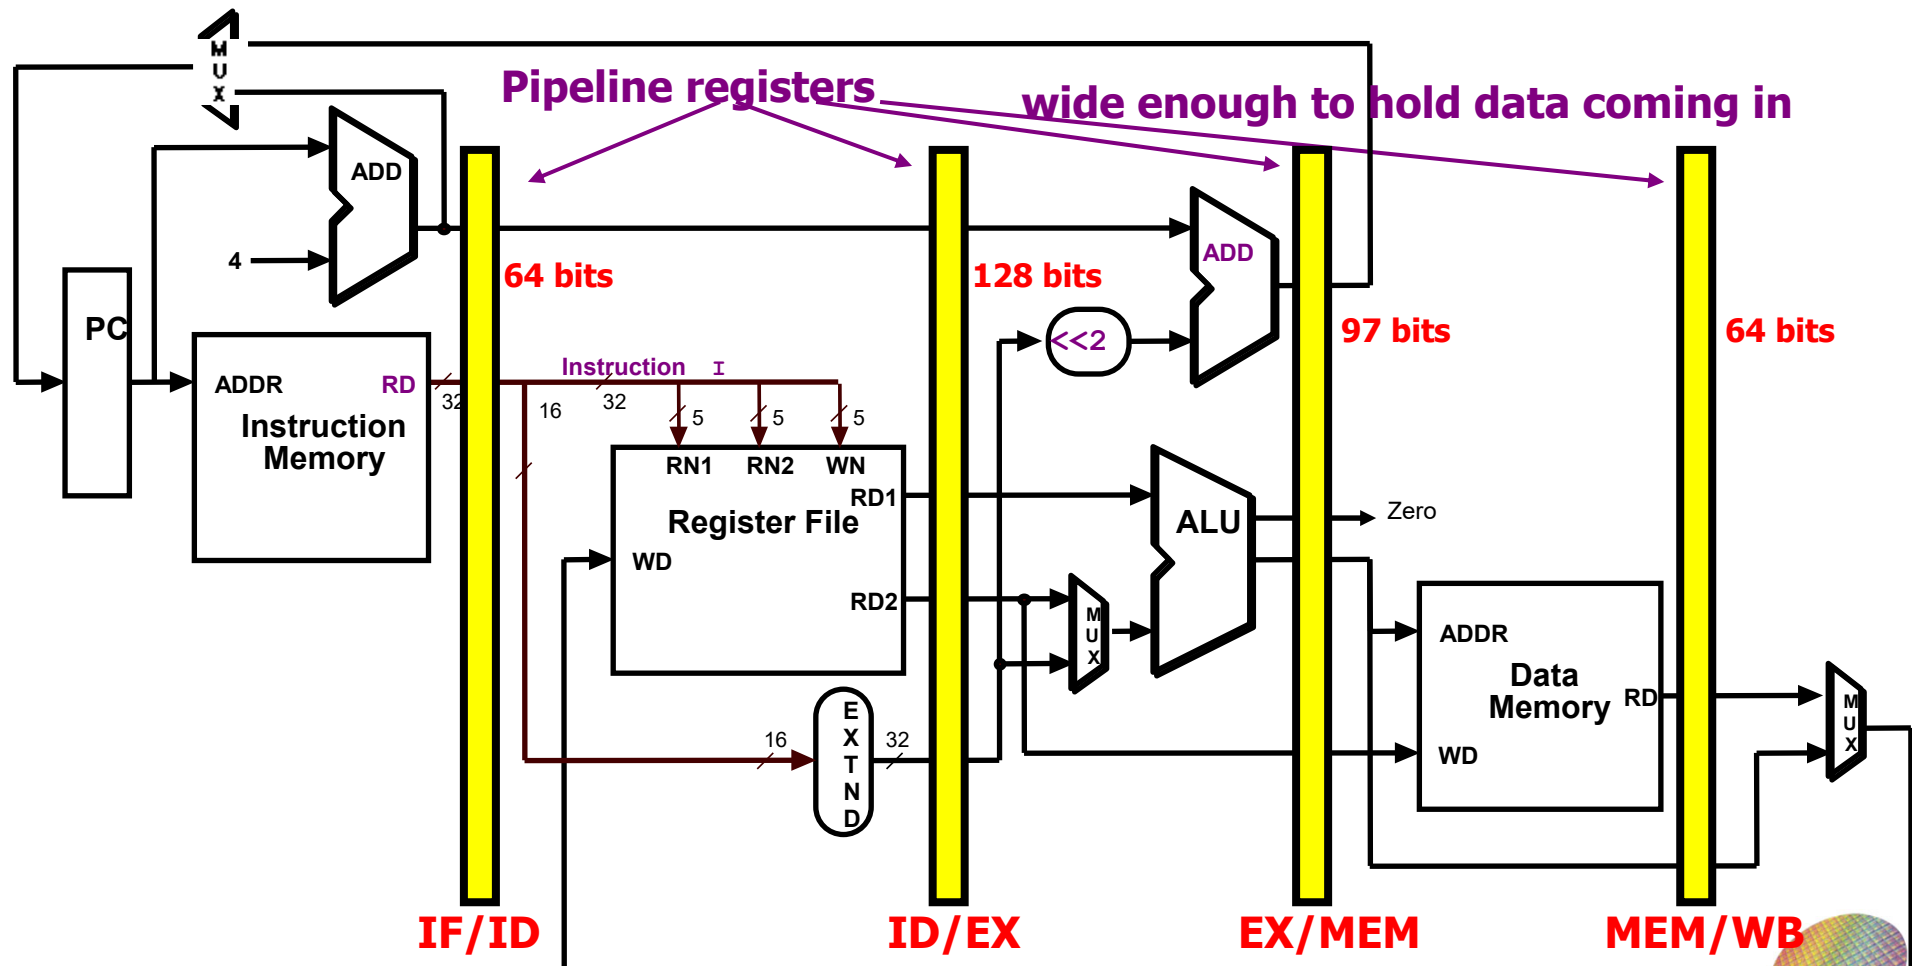
Not every instruction need 5 steps!

R-type need **4 steps**,
lw needs **5 steps**,
branch needs **3 steps**

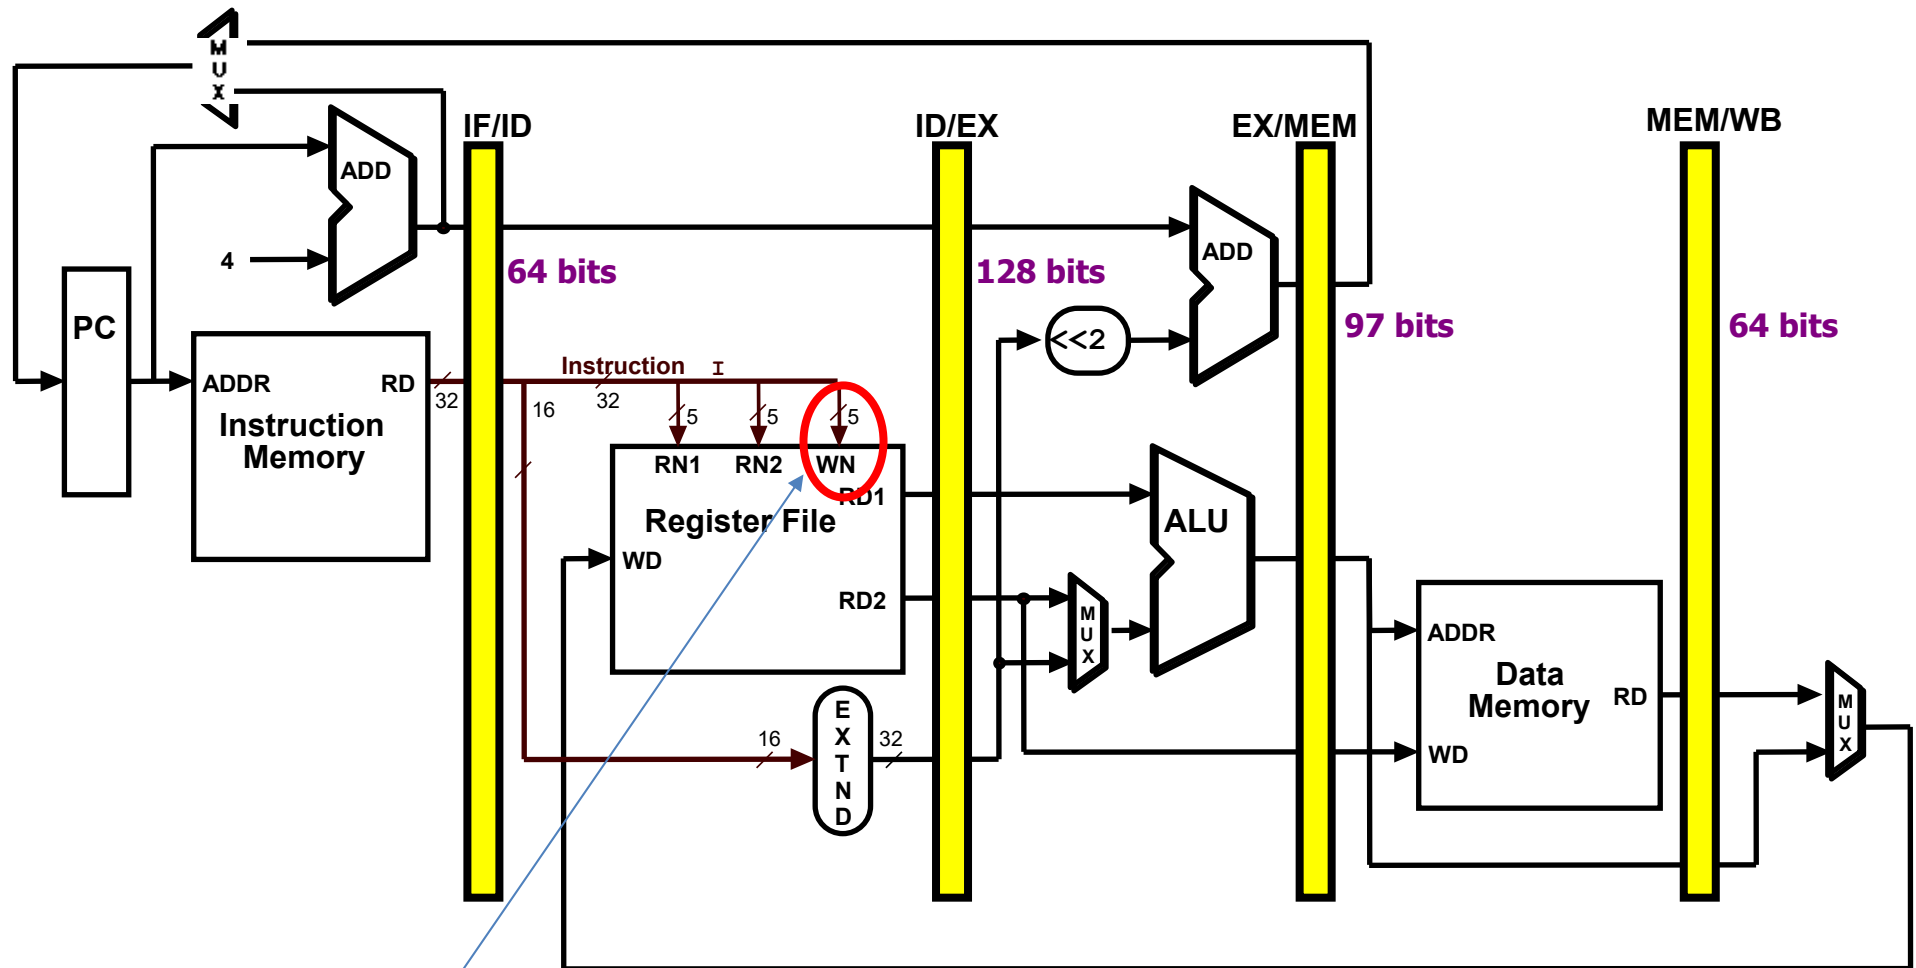


Pipelined Datapath


- Need registers between stages
 - To **hold** information produced in previous cycle
 - Data pass to next stage only during clock transition



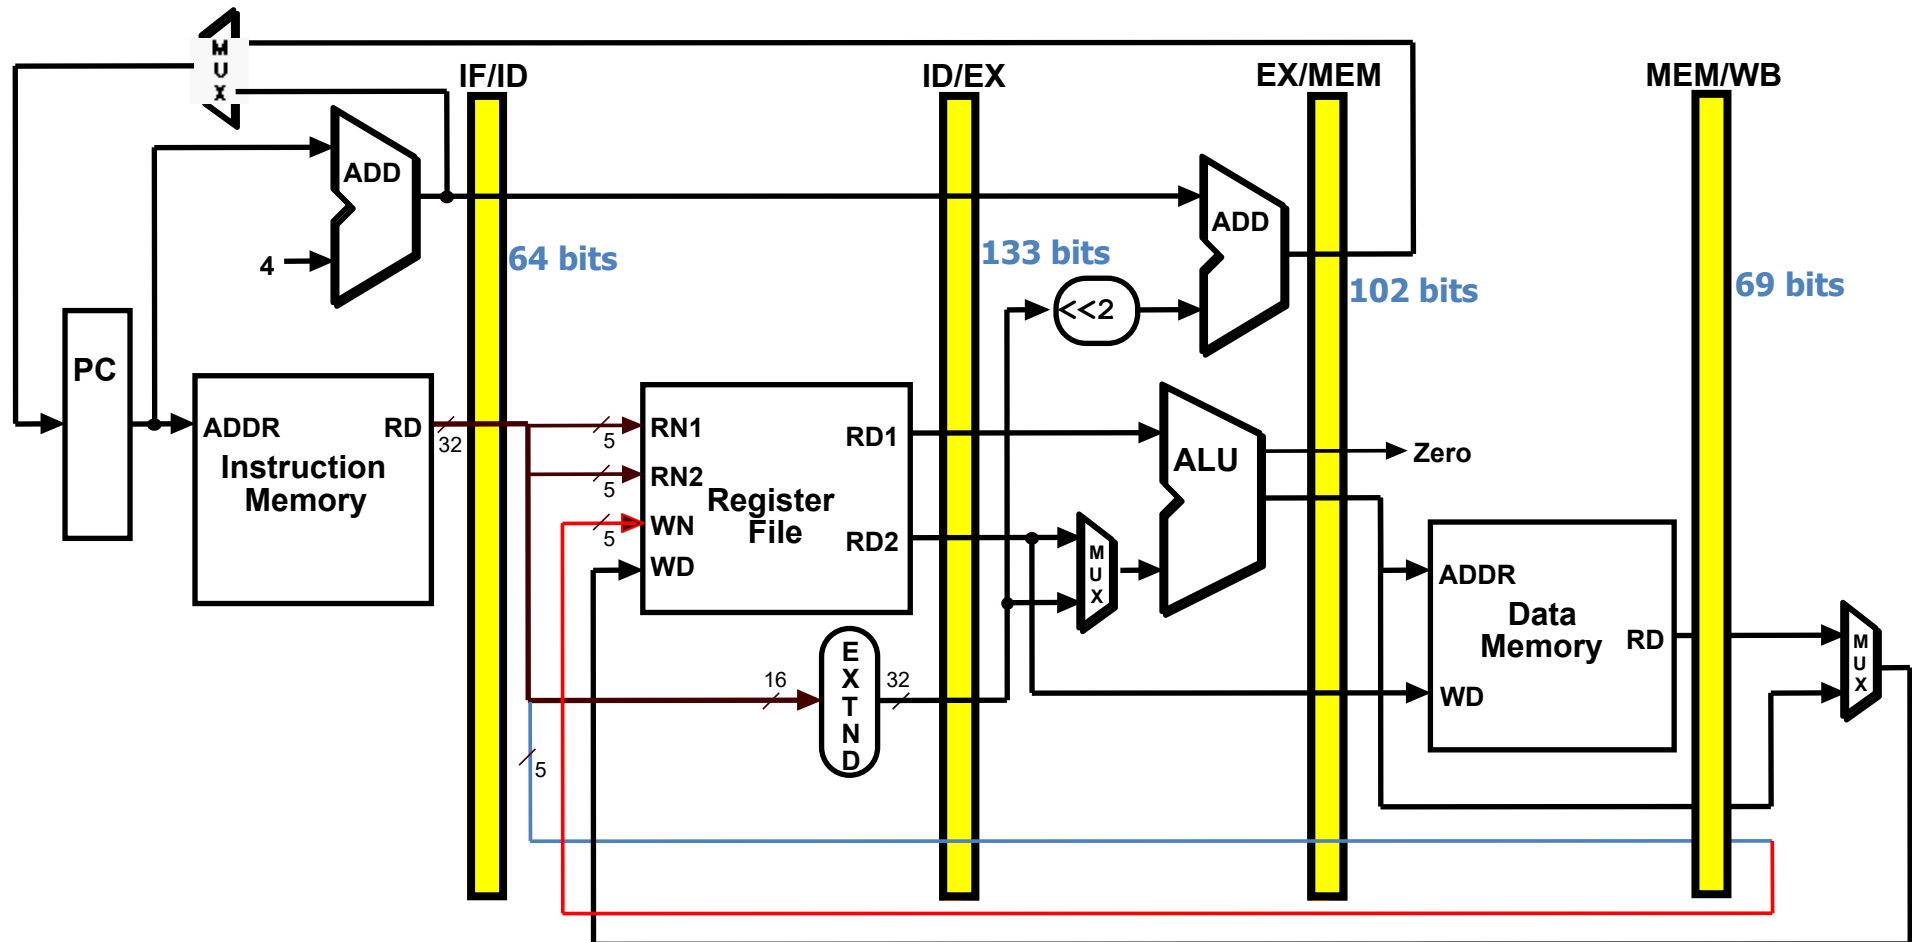
Bug in the Datapath




Write data from **Data Memory**, but Write register number comes from **another** instruction in the latter stage!



Corrected Datapath




Destination register number is also passed through ID/EX, EX/MEM and MEM/WB registers, which are now wider by 5 bits



Pipeline Operation

- Cycle-by-cycle flow of instructions through the pipelined datapath
 - “Single-clock-cycle” pipeline diagram
 - Shows pipeline usage in a single cycle
 - Highlight resources used
 - c.f. “multi-clock-cycle” diagram
 - Graph of operation over time
- We’ll look at “single-clock-cycle” diagrams for load & store



Pipelined Example


- Consider the following instruction sequence:


lw \$t0, 10 (\$t1)

sw \$t3, 20 (\$t4)

add \$t5, \$t6, \$t7


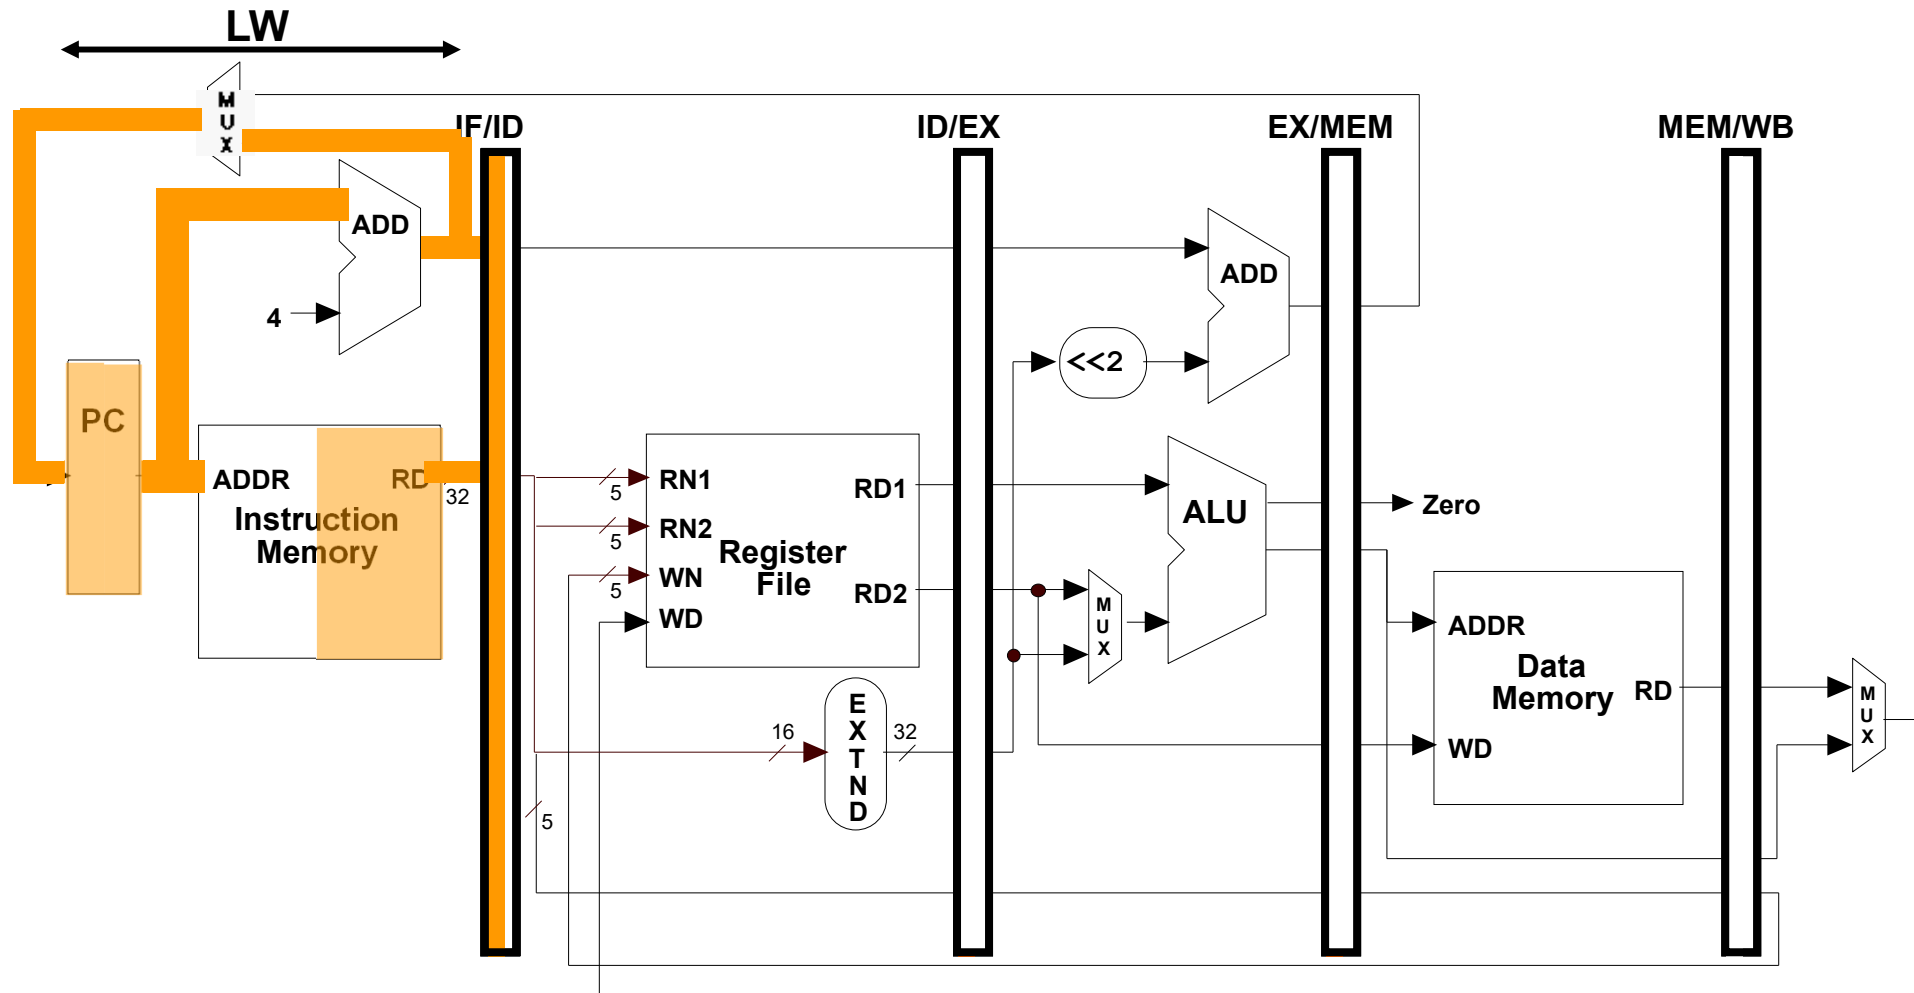
sub \$t8, \$t9, \$t10





Single-Clock-Cycle Diagram: Clock Cycle 1

lw \$t0, 10(\$t1)
sw \$t3, 20(\$t4)
add \$t5, \$t6, \$t7
sub \$t8, \$t9, \$t10




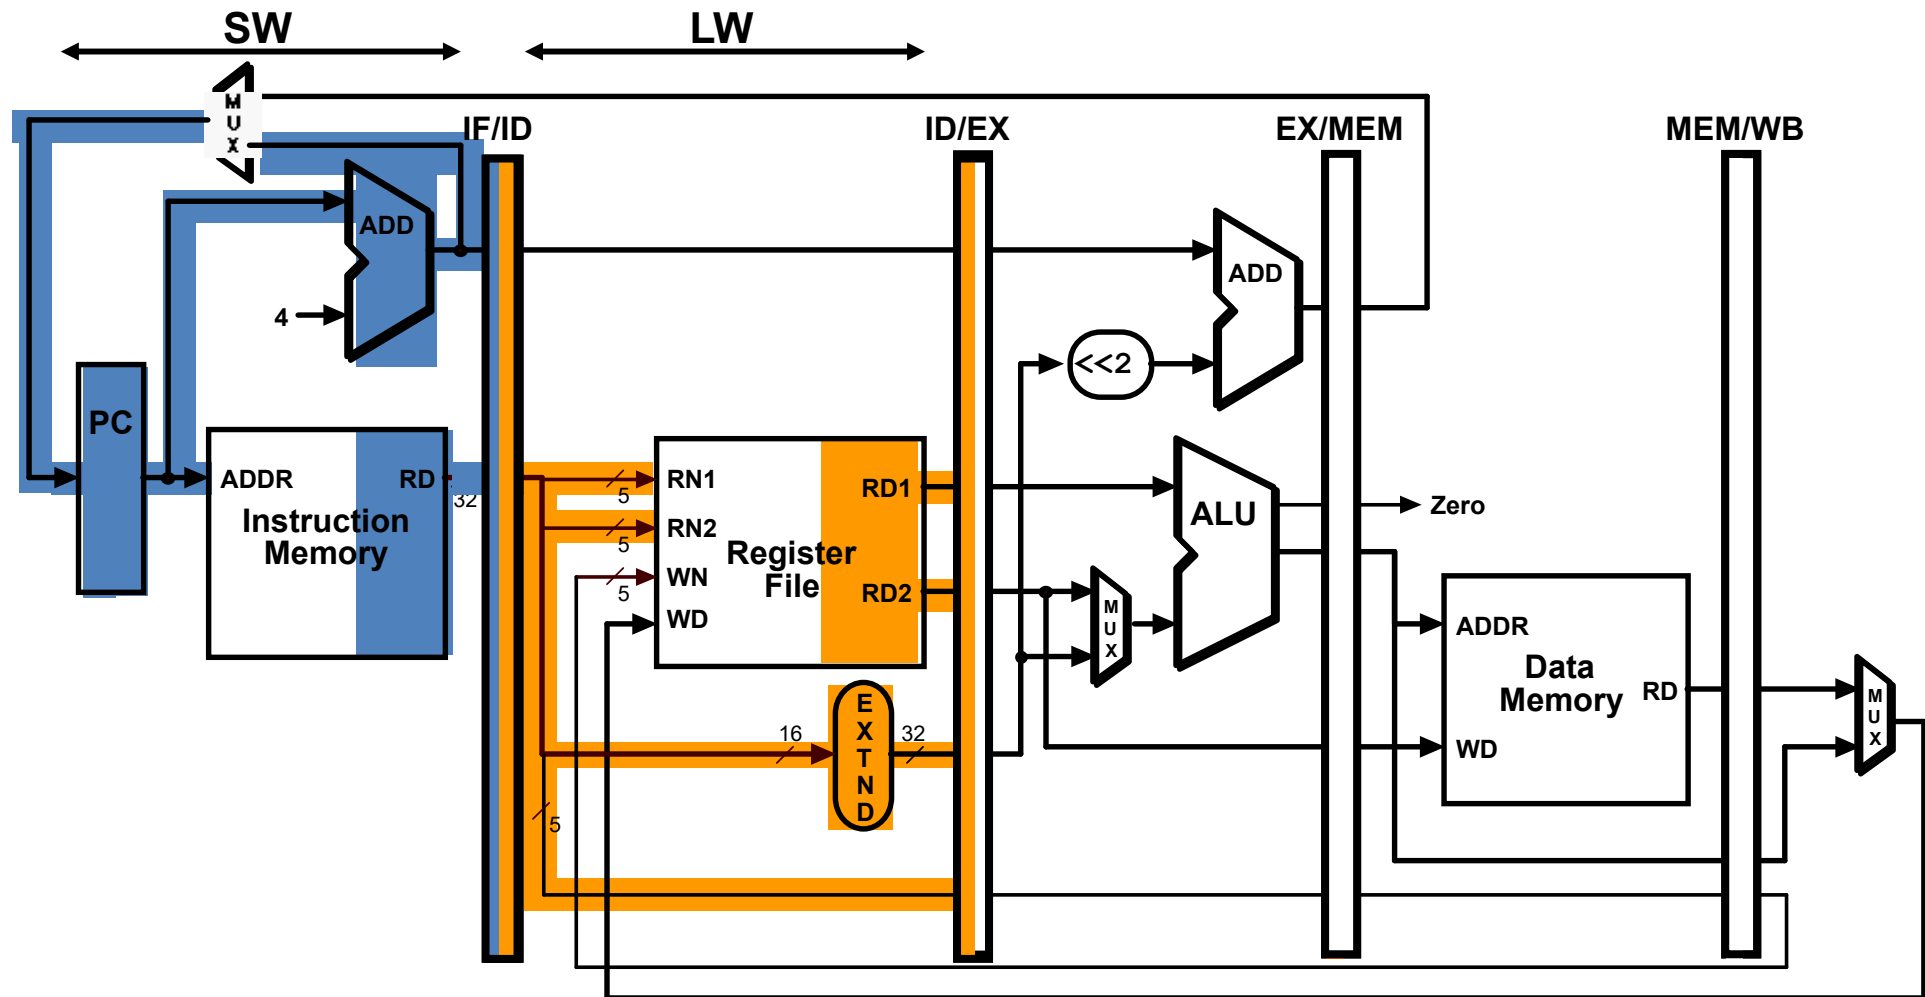
Single-Clock-Cycle Diagram: Clock Cycle 2


```
lw $t0, 10($t1)
```

```
sw $t3, 20($t4)
```

```
add $t5, $t6, $t7
```


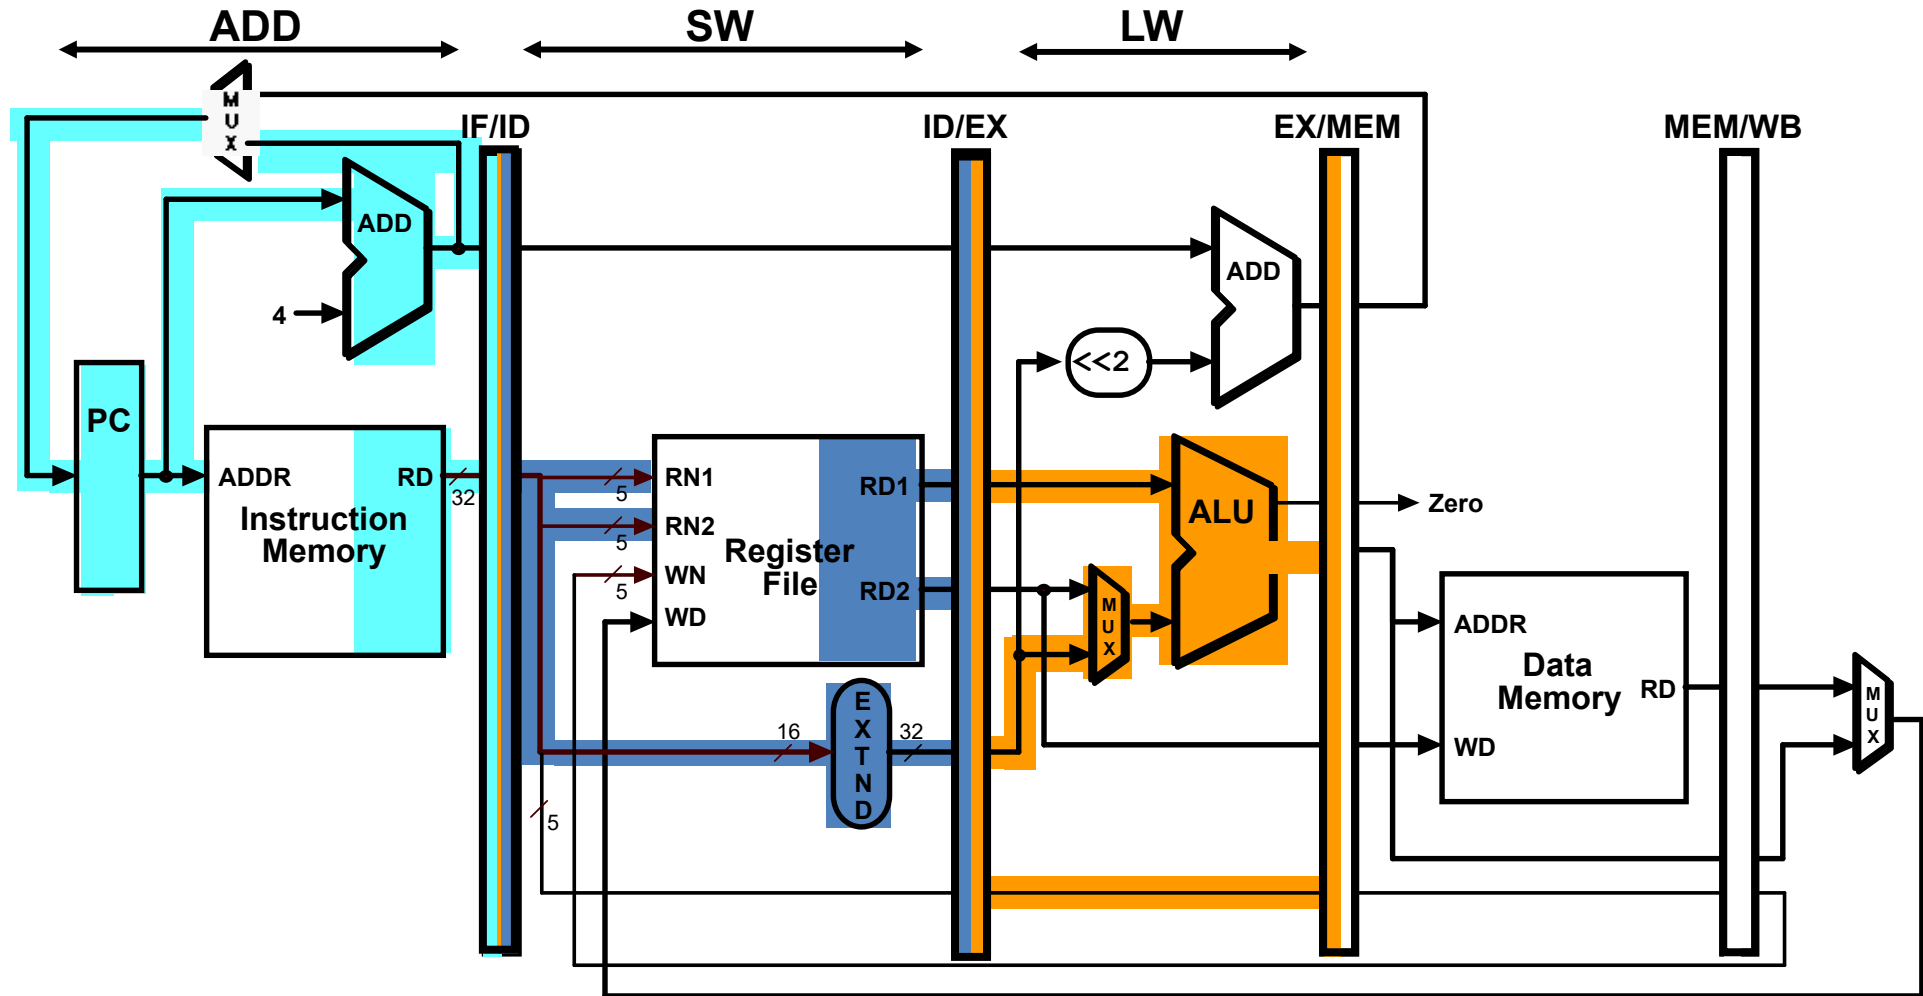
```
sub $t8, $t9, $t10
```






Single-Clock-Cycle Diagram: Clock Cycle 3


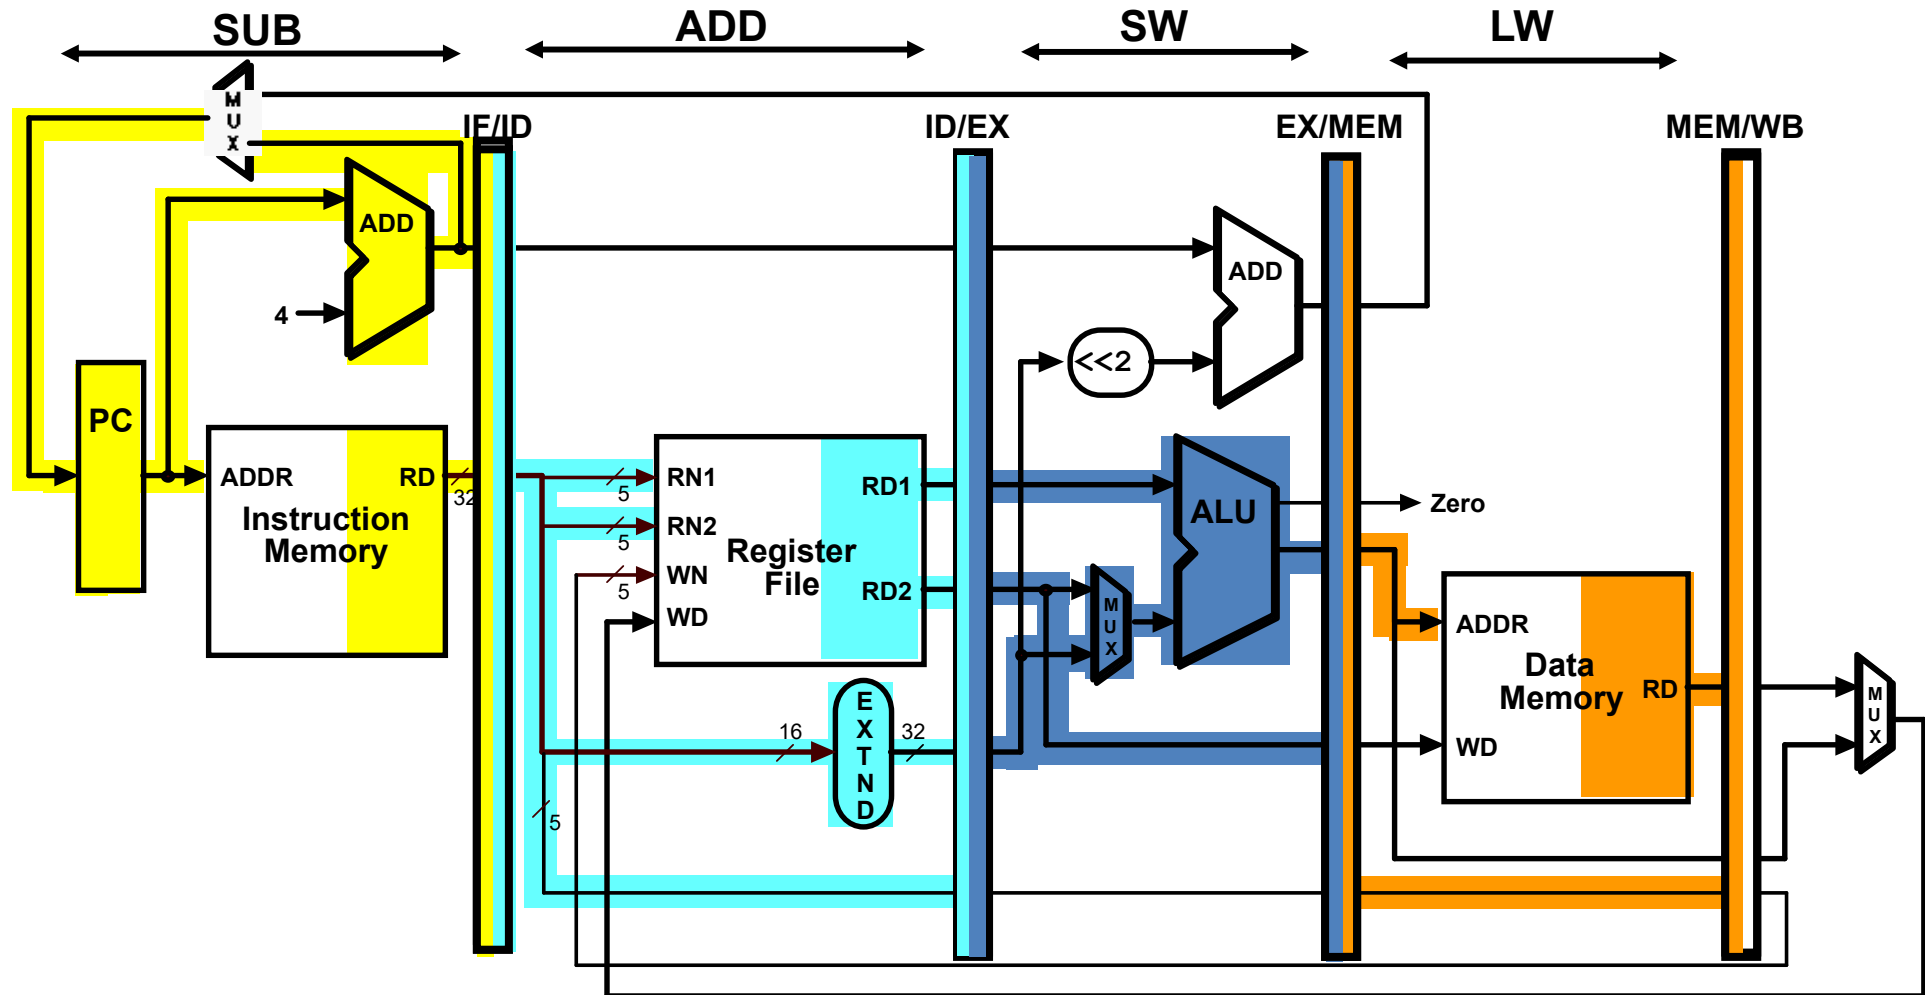
lw \$t0, 10(\$t1)
sw \$t3, 20(\$t4)
add \$t5, \$t6, \$t7
sub \$t8, \$t9, \$t10






Single-Clock-Cycle Diagram: Clock Cycle 4


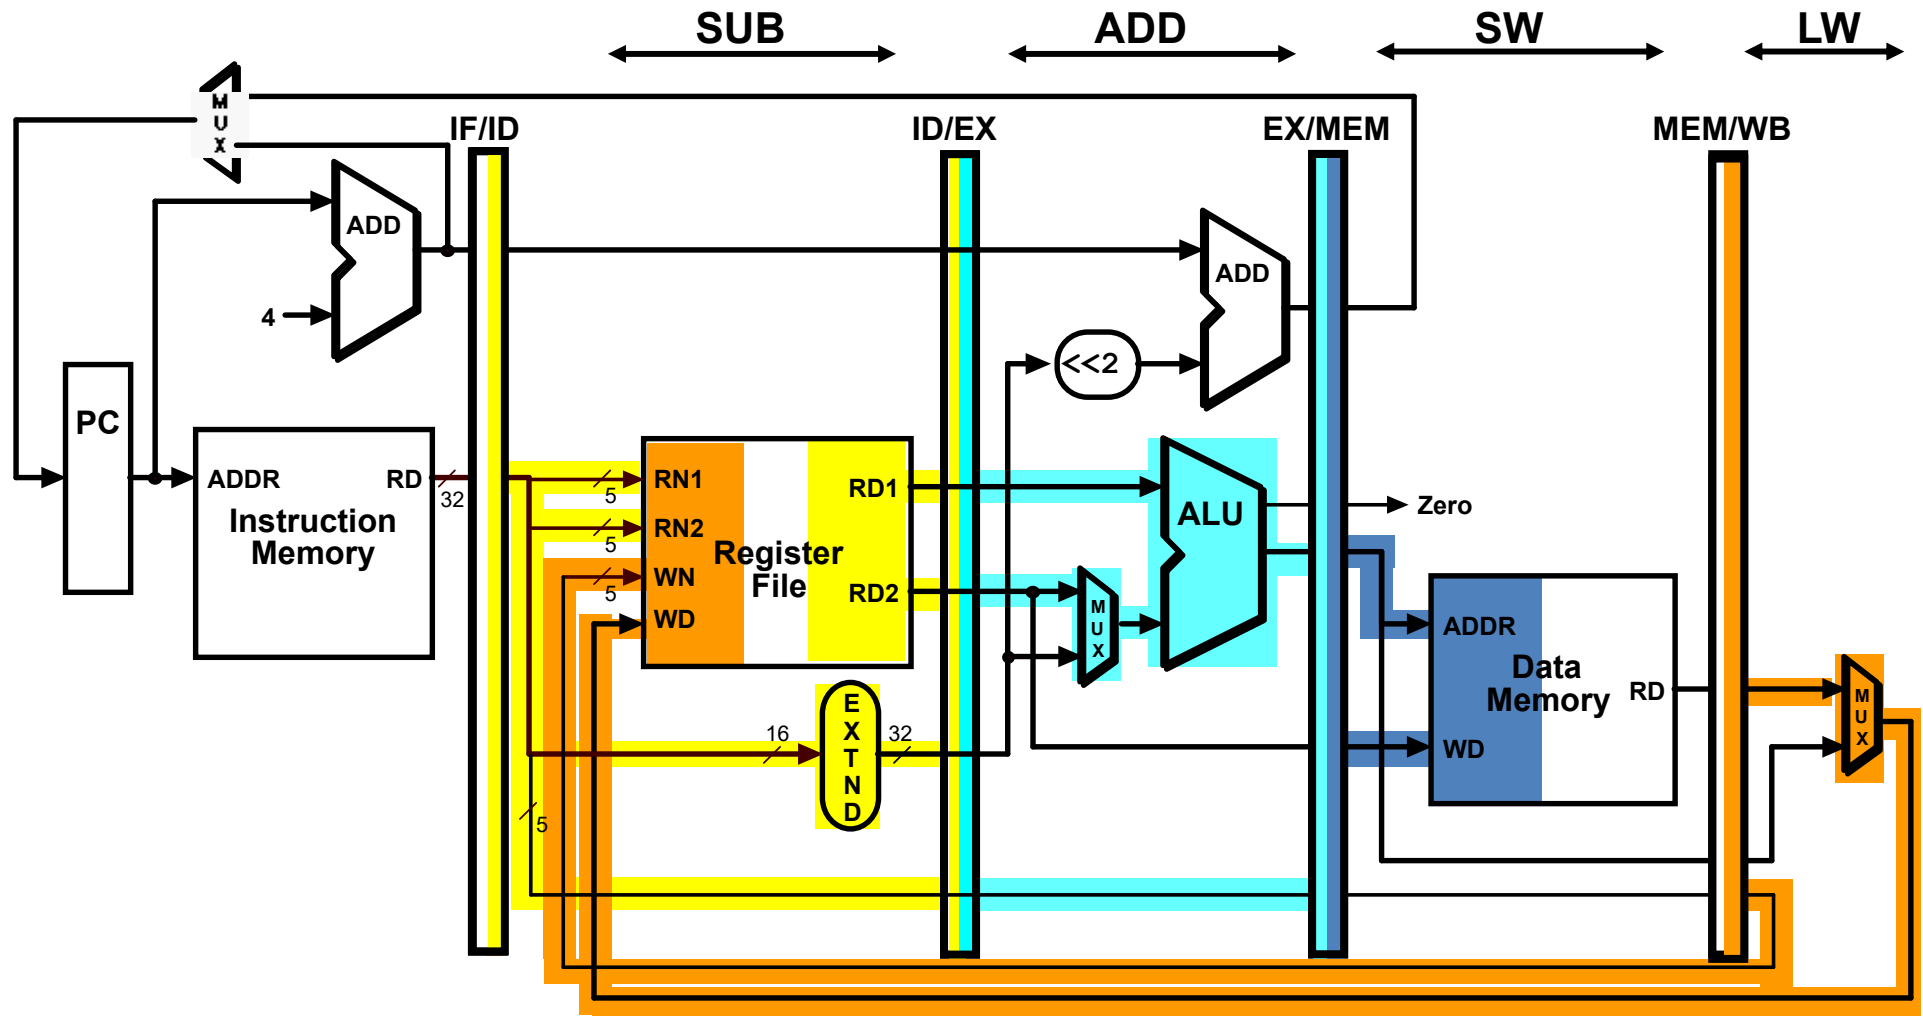
lw \$t0, 10(\$t1)
sw \$t3, 20(\$t4)
add \$t5, \$t6, \$t7
sub \$t8, \$t9, \$t10





Single-Clock-Cycle Diagram: Clock Cycle 5

lw \$t0, 10(\$t1)
sw \$t3, 20(\$t4)
add \$t5, \$t6, \$t7
sub \$t8, \$t9, \$t10




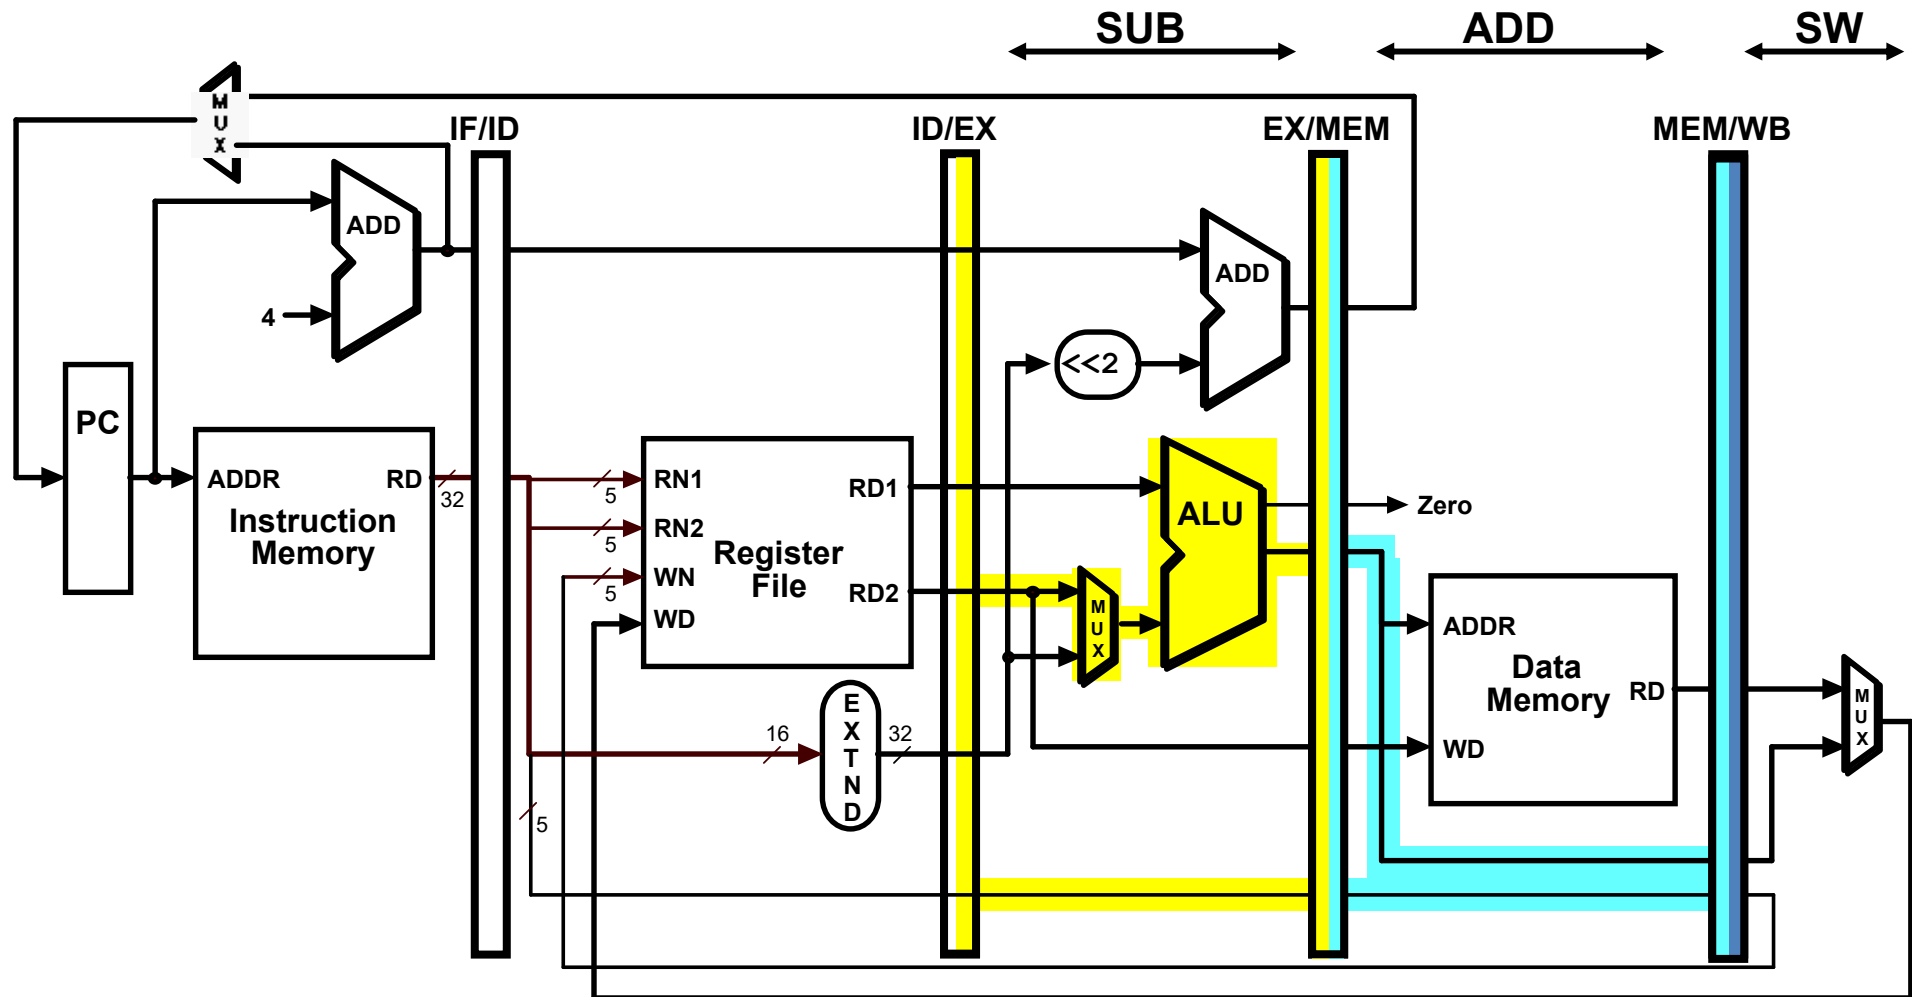
Single-Clock-Cycle Diagram: Clock Cycle 6


```
lw $t0, 10($t1)
```

```
sw $t3, 20($t4)
```

```
add $t5, $t6, $t7
```


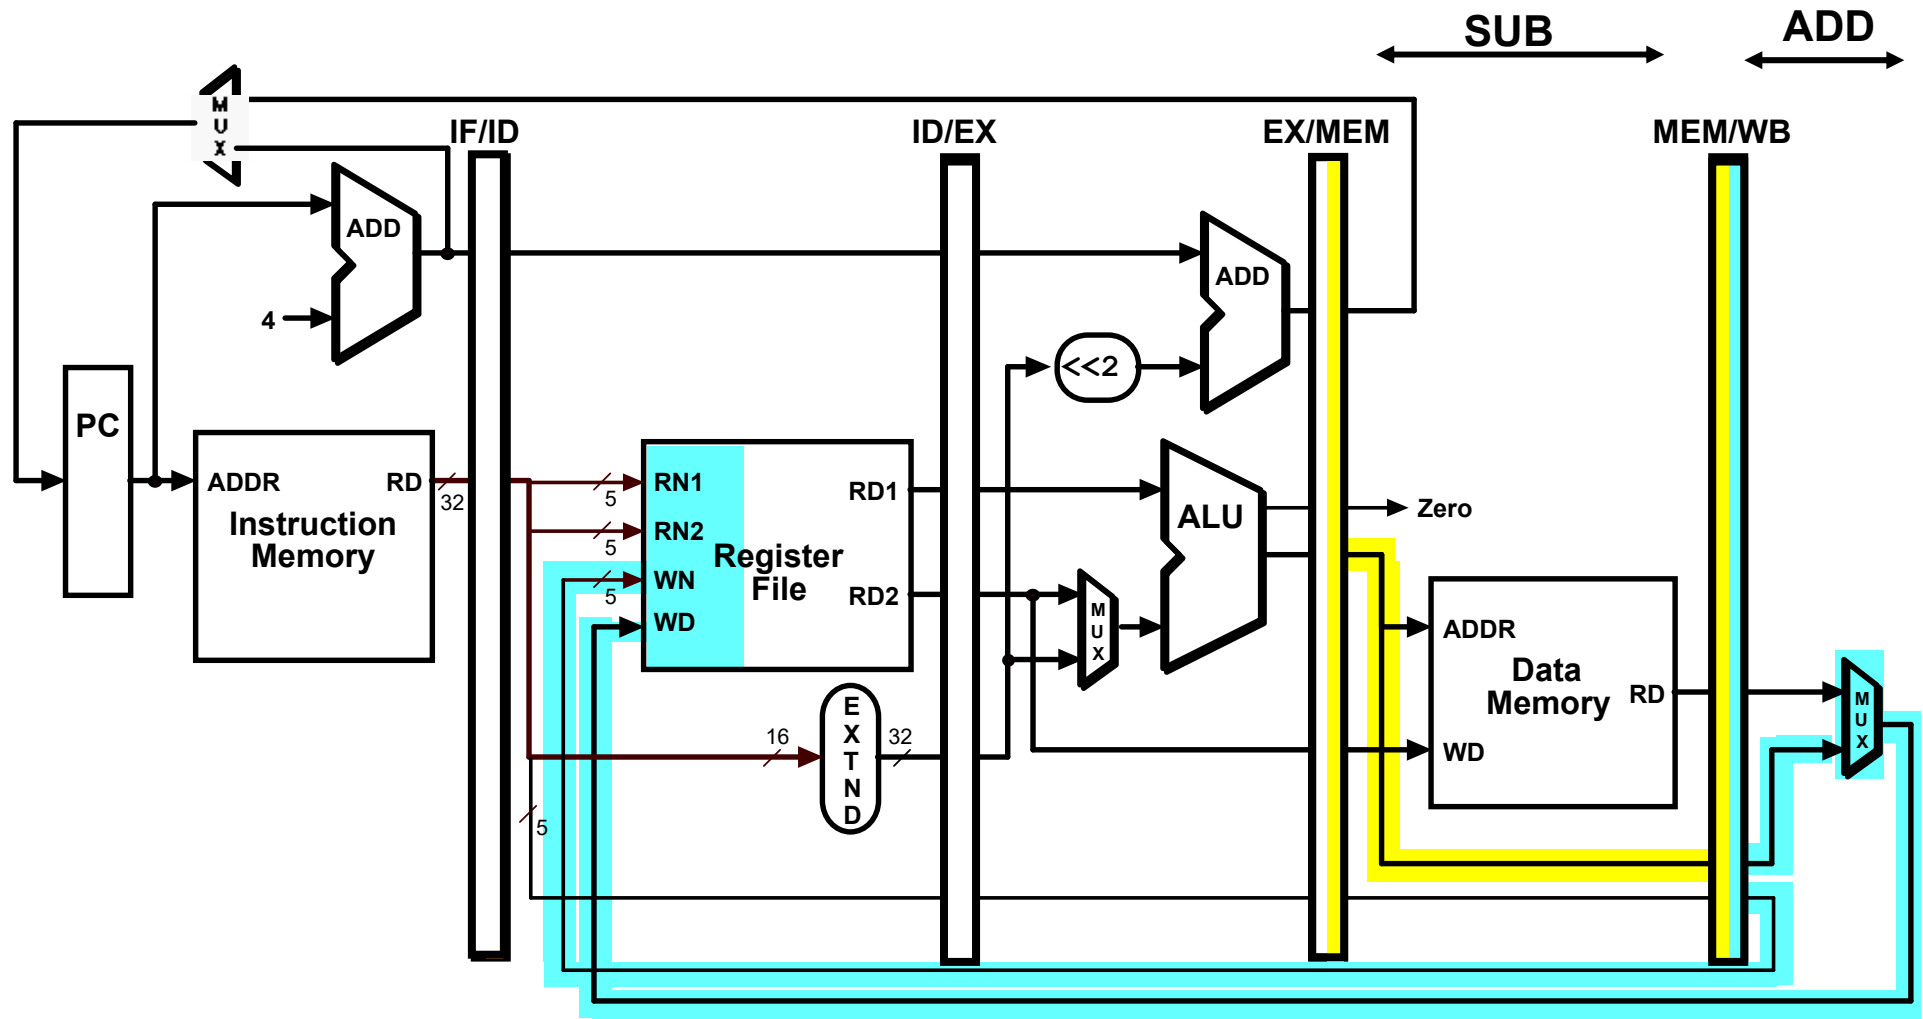
```
sub $t8, $t9, $t10
```






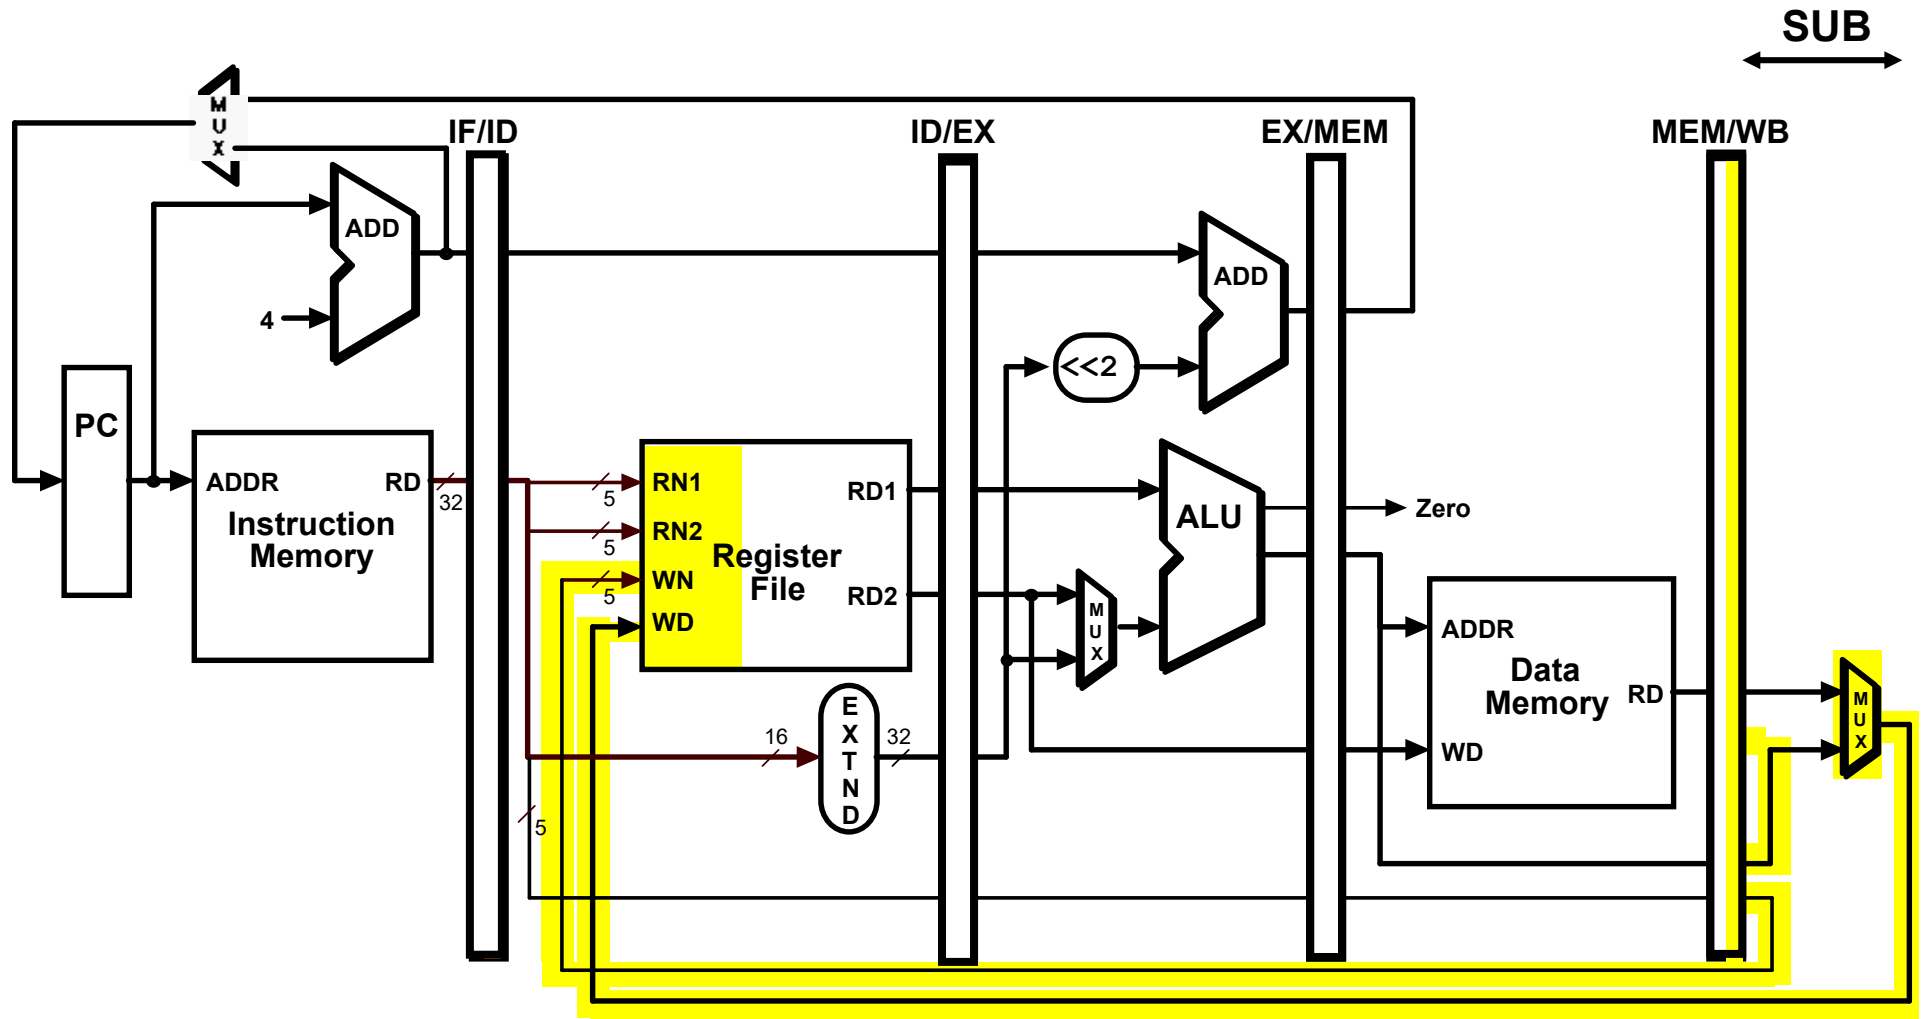
Single-Clock-Cycle Diagram: Clock Cycle 7

lw \$t0, 10(\$t1)
sw \$t3, 20(\$t4)
add \$t5, \$t6, \$t7
sub \$t8, \$t9, \$t10




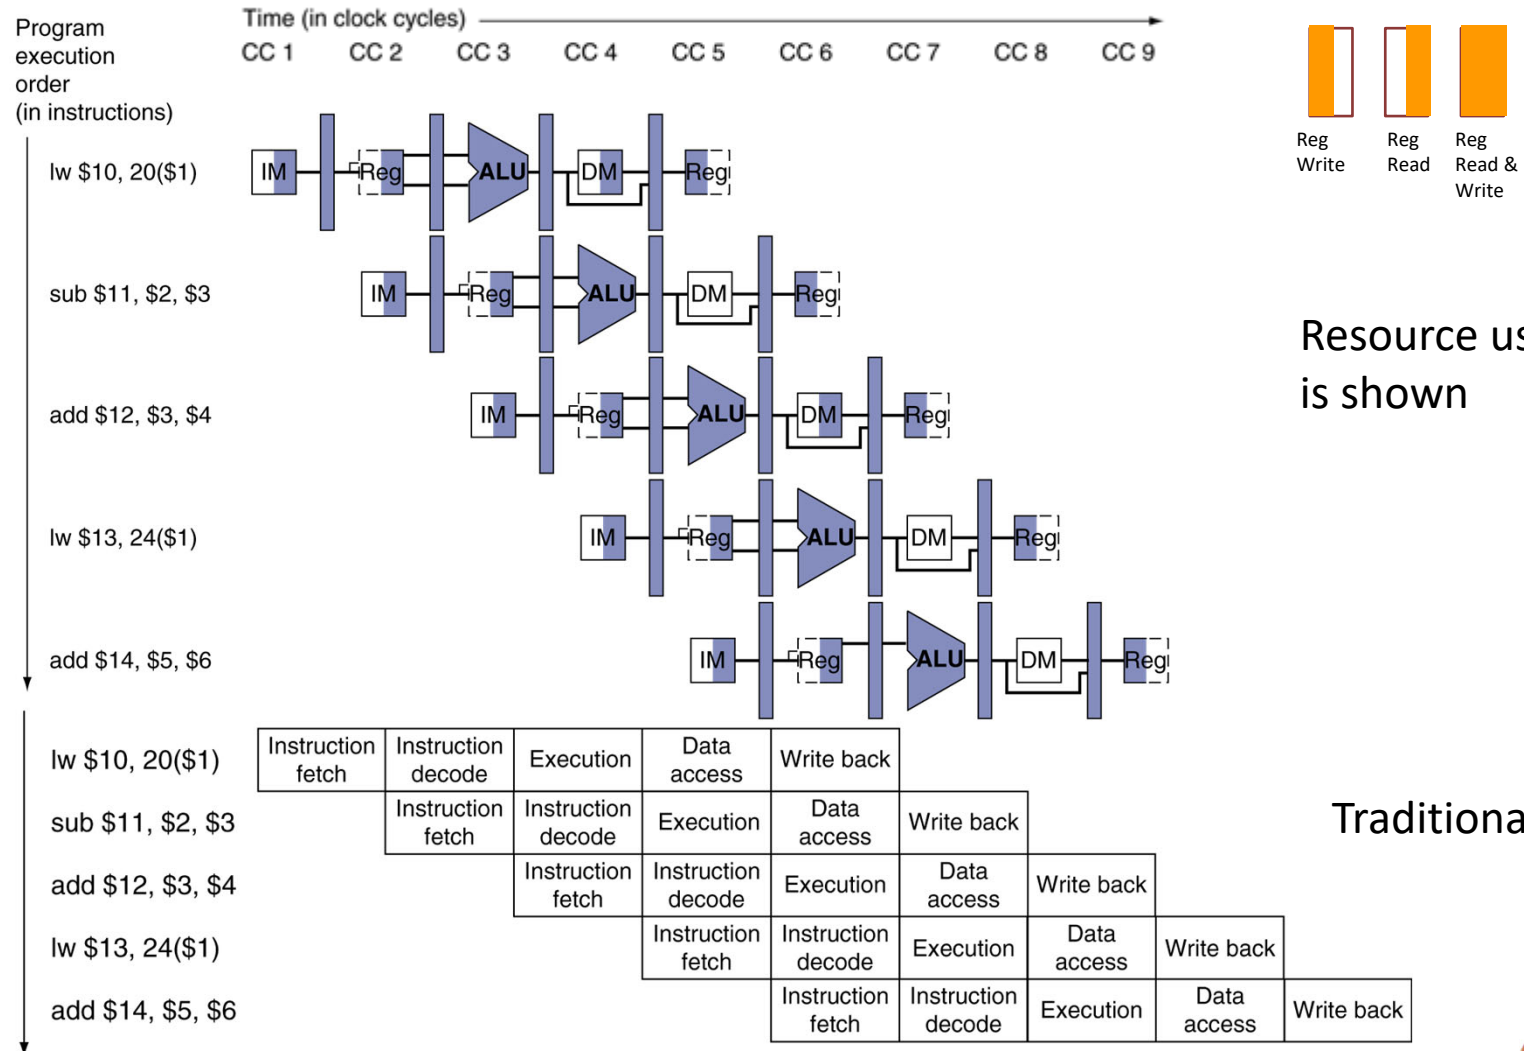
Single-Clock-Cycle Diagram: Clock Cycle 8

lw \$t0, 10(\$t1)
sw \$t3, 20(\$t4)
add \$t5, \$t6, \$t7
sub \$t8, \$t9, \$t10



Multi-Cycle Pipeline Diagram

- Two ways to represent pipeline






成功大學

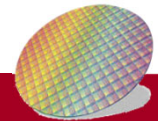
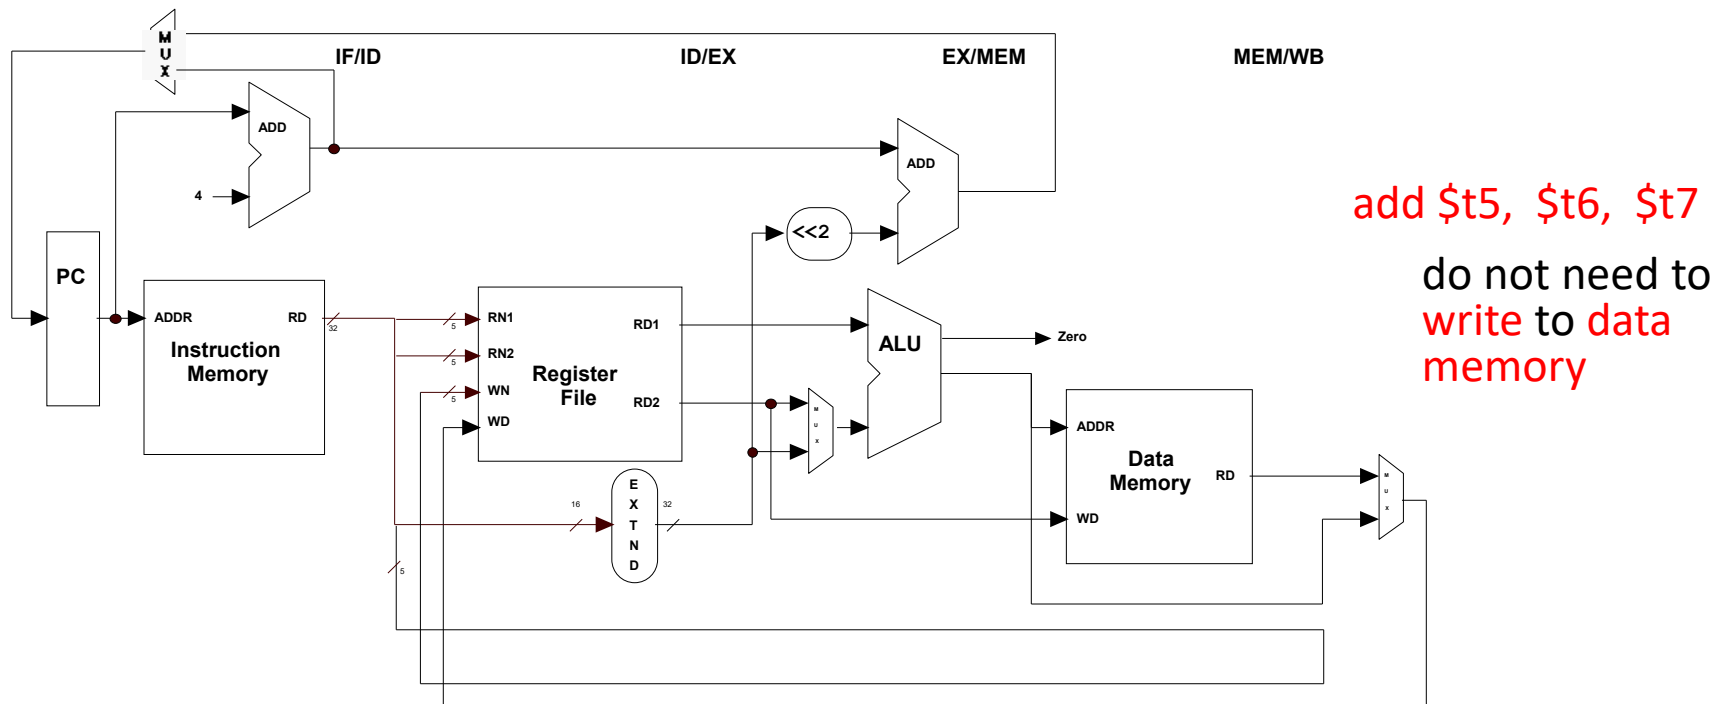
National Cheng Kung University

Why the execution of an R-type instruction
uses 5 stages?



Observation


- R-type instruction, do not need to **write to data memory**
 - Only IF, ID, EXE, WB stages are used
- why the execution of an R-type instruction uses **5 stages?**




Supposed R-type Instructions only use 4 stages in pipeline



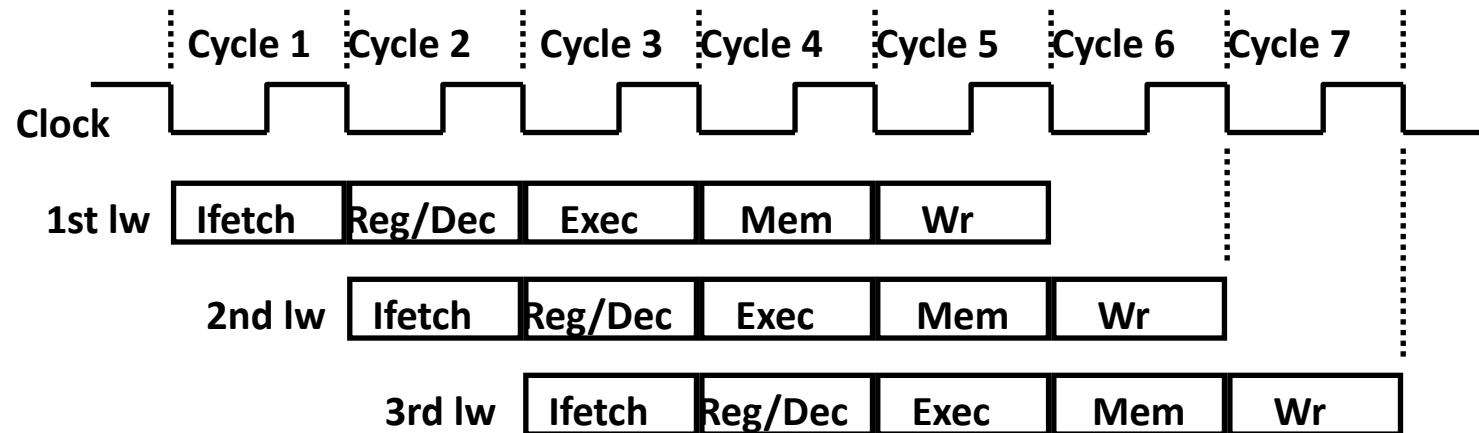
成功大學




- **IF**: fetch the instruction from the Instruction Memory
- **ID**: registers fetch and instruction decode
- **EX**: ALU operates on the two register operands
- **WB**: write ALU output back to the register file



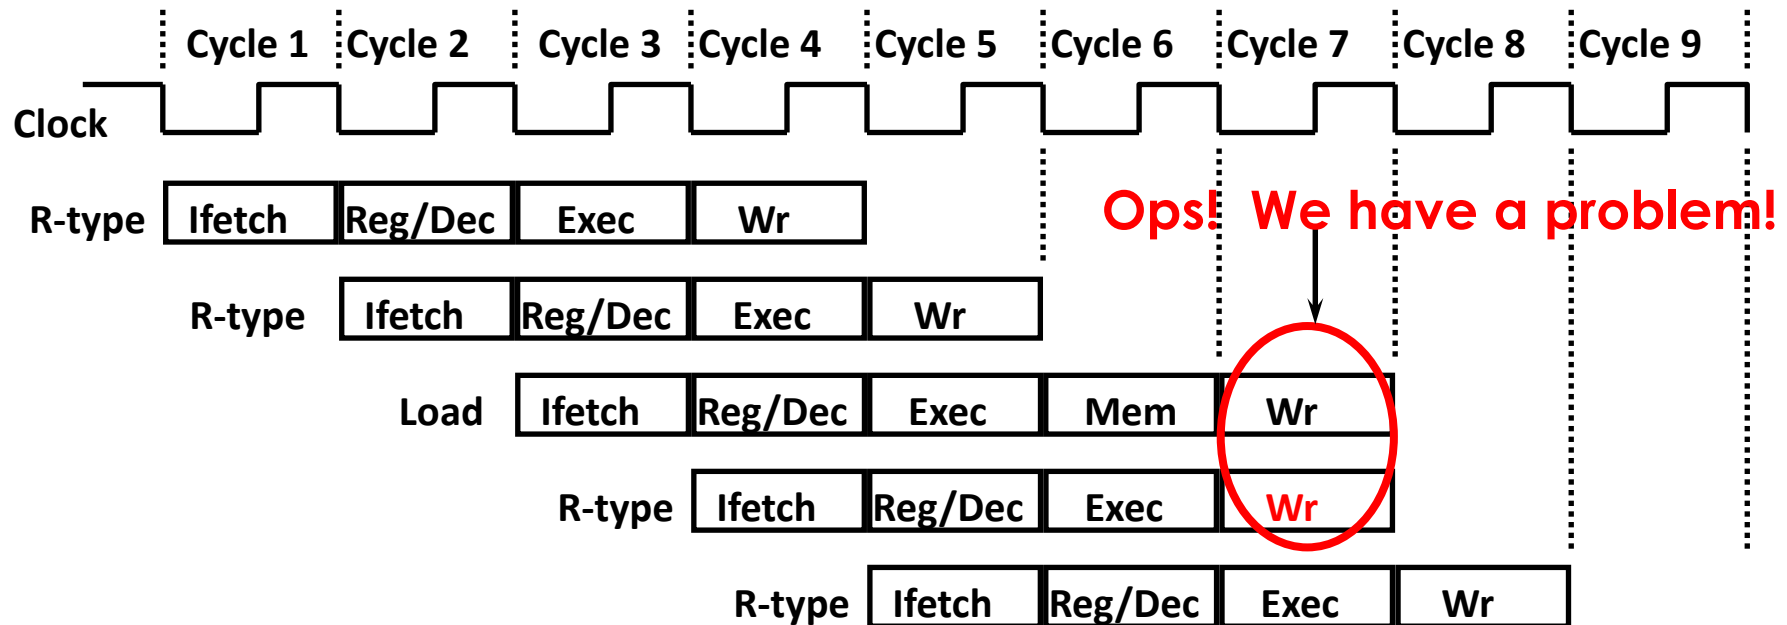
And Load has 5 stages




- 5 functional units in the pipeline datapath are:
 - **Instruction Memory** for the Fetch stage
 - Register File's **Read ports** (busA and busB) for the Reg/Dec stage
 - **ALU** for the Exec stage
 - **Data Memory** for the MEM stage
 - Register File's **Write port** (busW) for the WB stage



Pipelining with R-type and load


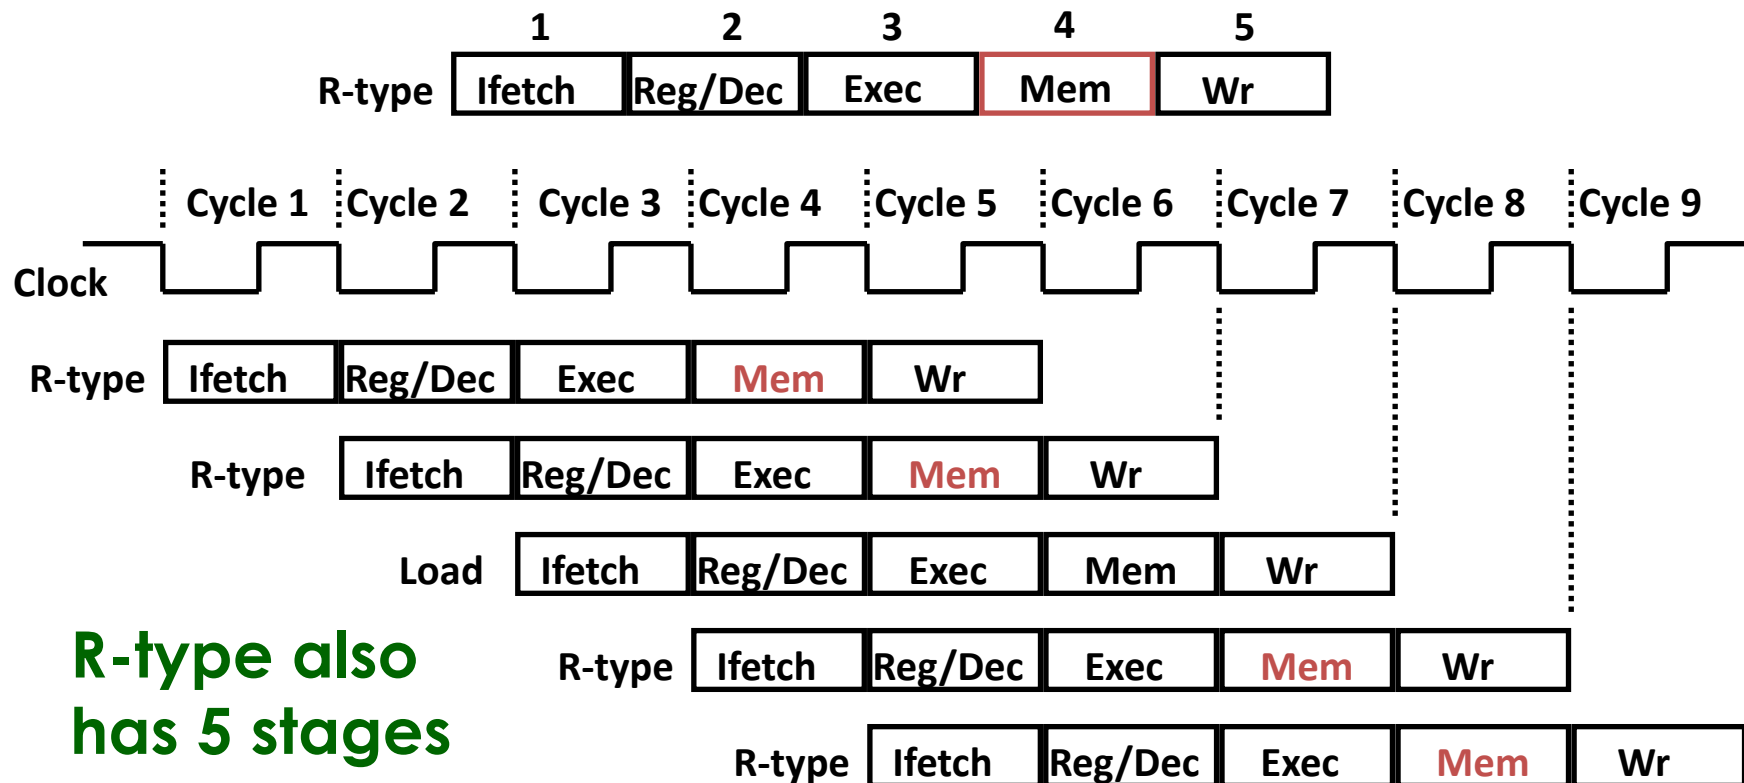


- We have a *structural hazard*:
 - Two instructions try to write to the register file at the same time!
 - Only one write port



Solution: Delay R-type's Write

- Delay R-type's register write by one cycle:
 - R-type also use Reg File's write port at Stage 5
 - MEM is a **NOP** stage: nothing is being done.





SW

- Similar situation applies to **SW**, originally only 4 stages are needed
 - IF: fetch the instruction from the Instruction Memory
 - ID: registers fetch and instruction decode
 - EX: calculate the memory address
 - MEM: write the data into the Data Memory

Add an extra stage:

- WB: **NOP (No operation)**



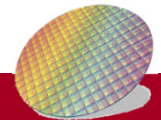

BEQ

Similar situation applies to **BEQ** instruction, originally only 3 stages are needed

- IF: fetch the instruction from the Instruction Memory
- ID: registers fetch and instruction decode
- EX: compare the two register operand, select correct branch target address, and latch into PC

Add two extra stages:

- **MEM: NOP**
- **WB: NOP**





成功大學

National Cheng Kung University

Backup slides

